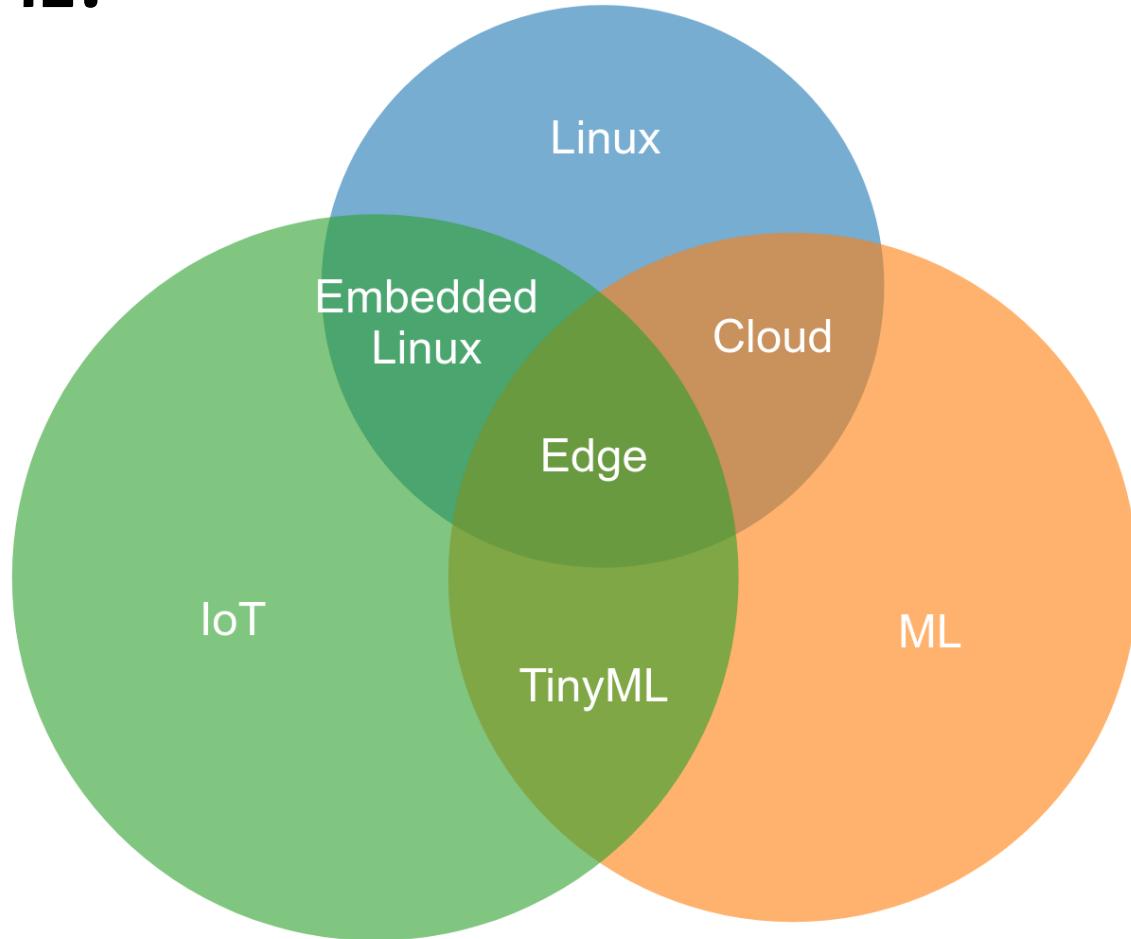


Distributed ML on Unikernel for IoT

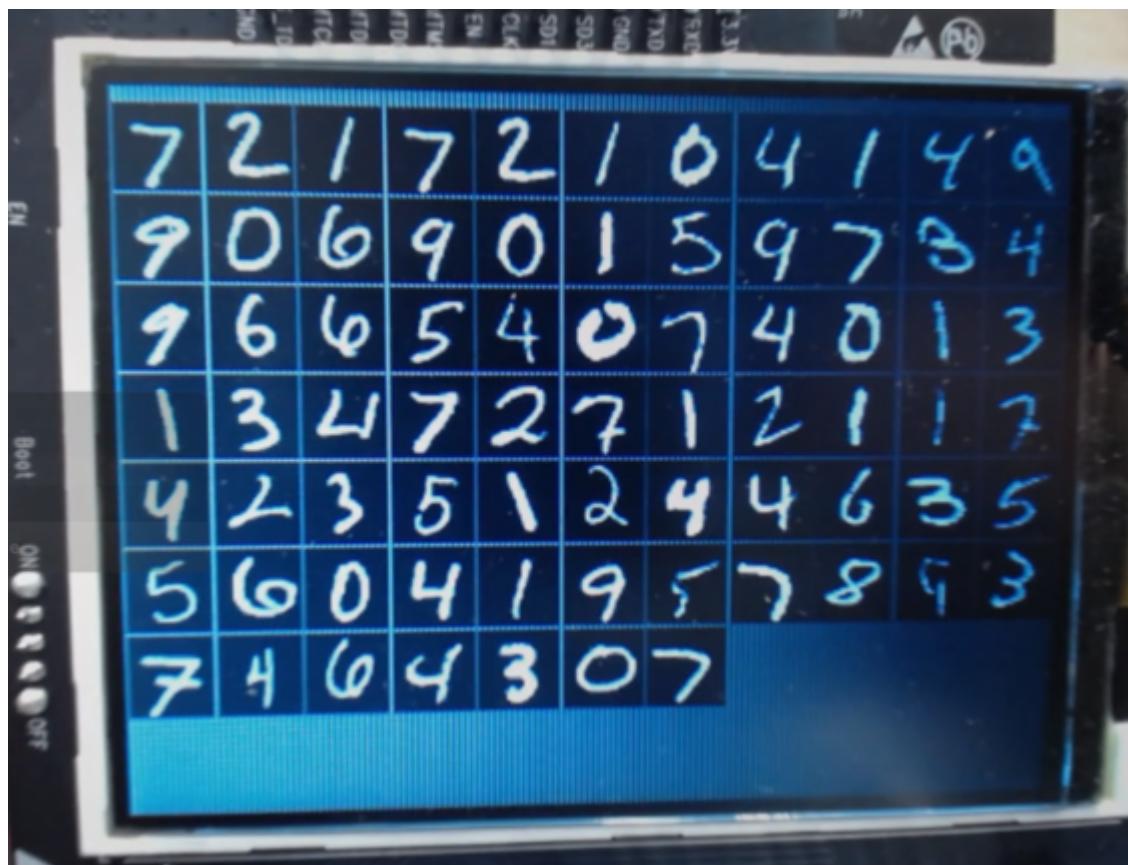
TinyML as-a-Service

Hiroshi Doyu <hiroshi.doyu@ericsson.com>

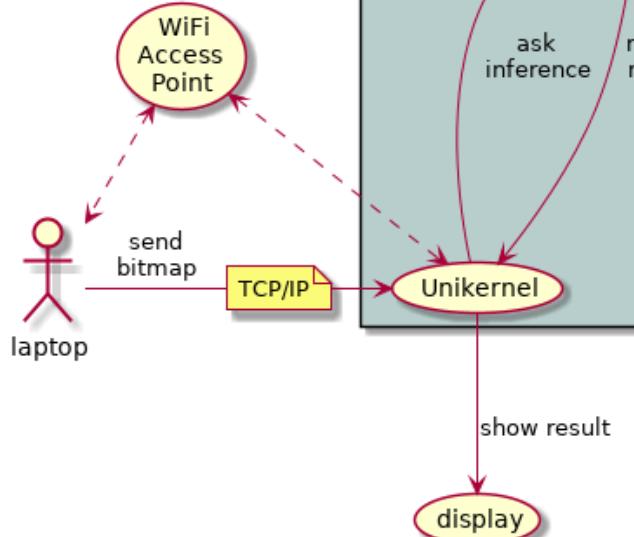
TinyML?



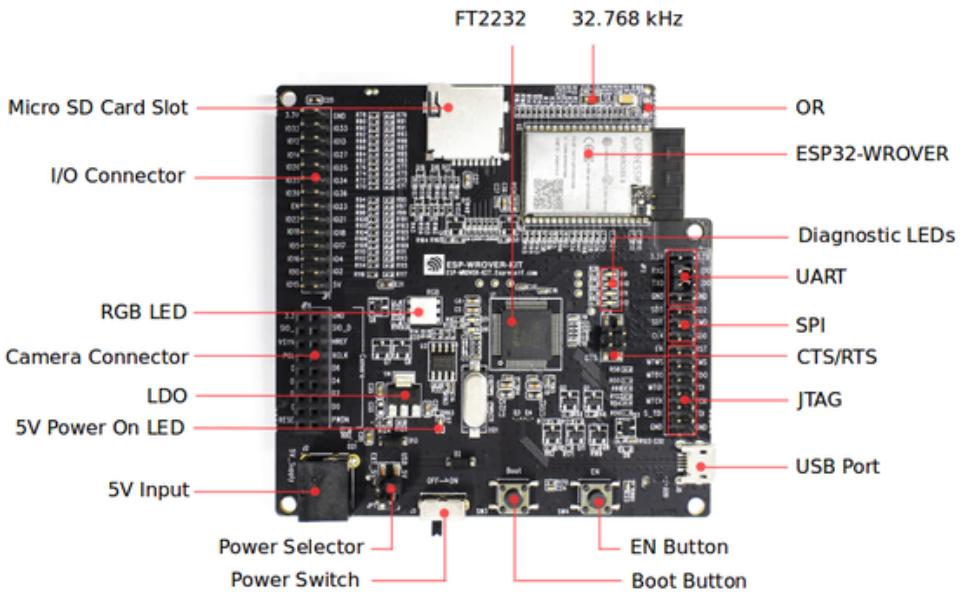
Demo

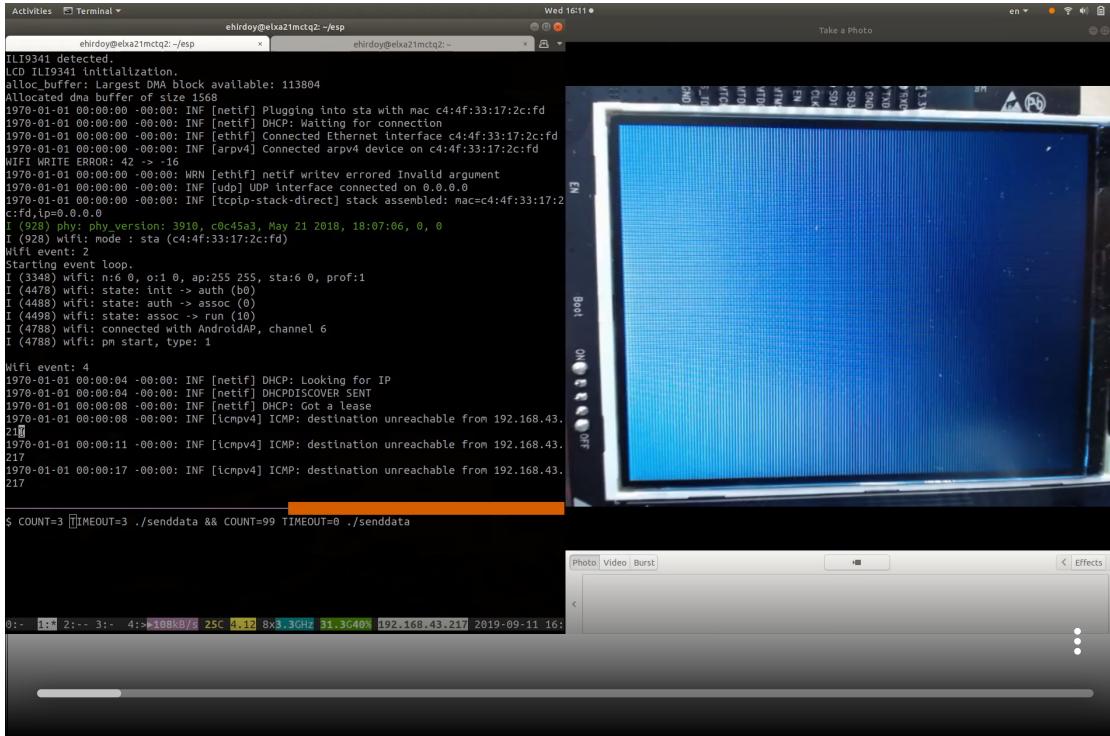


ML inference on MCU



MCU: ESP32
ARCH: Xtensa 32bit
RAM: 520KB SRAM
ROM: 4MB FLASH
160MHz / 600 DMIPS



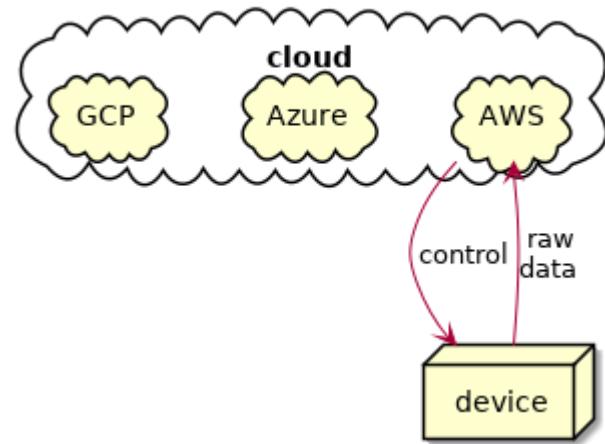


MNIST on ESP32

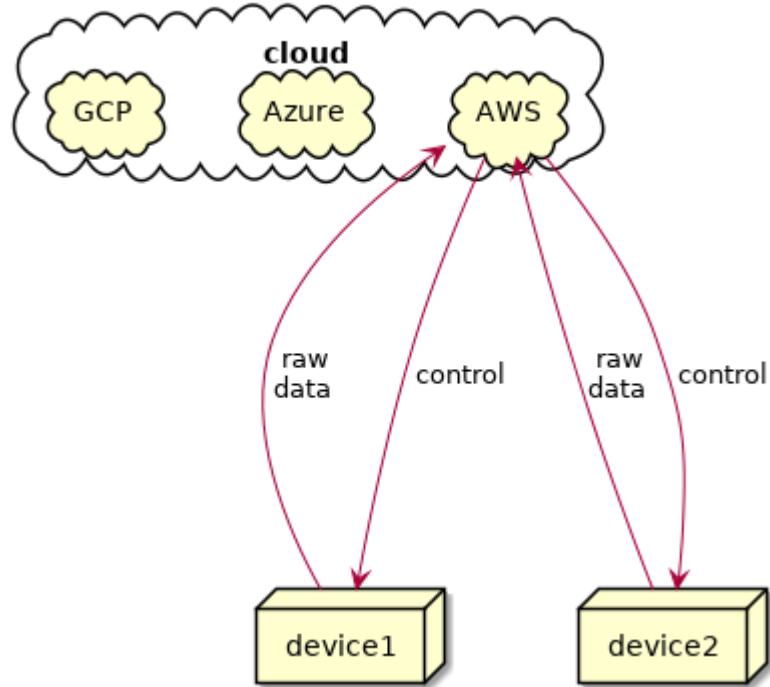
Outline

1. ~~Demo~~
2. **Problems**
 - **Edge Computing**
 - Web vs Embedded
 - ML environment
3. Proposal
4. Three Enablers
5. PoC
6. Conclusion

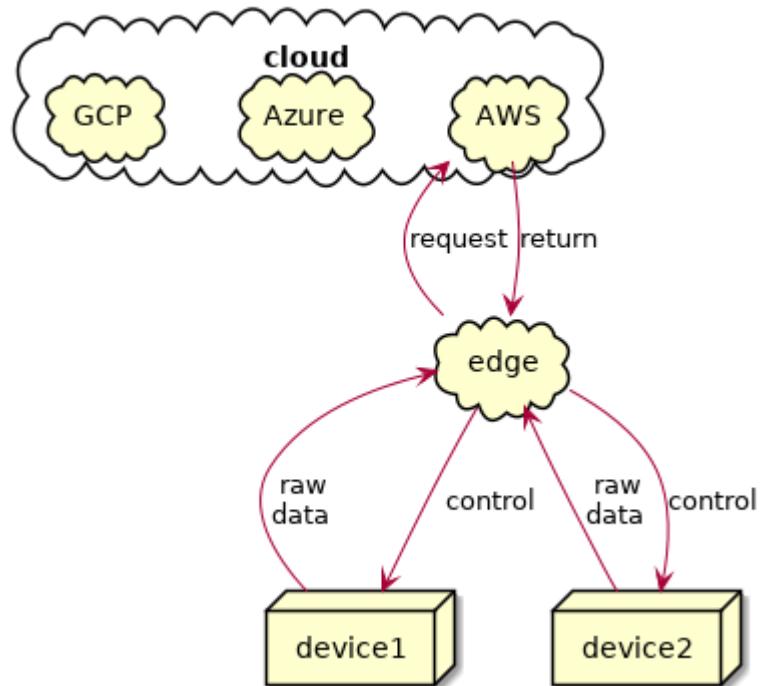
Traditionally



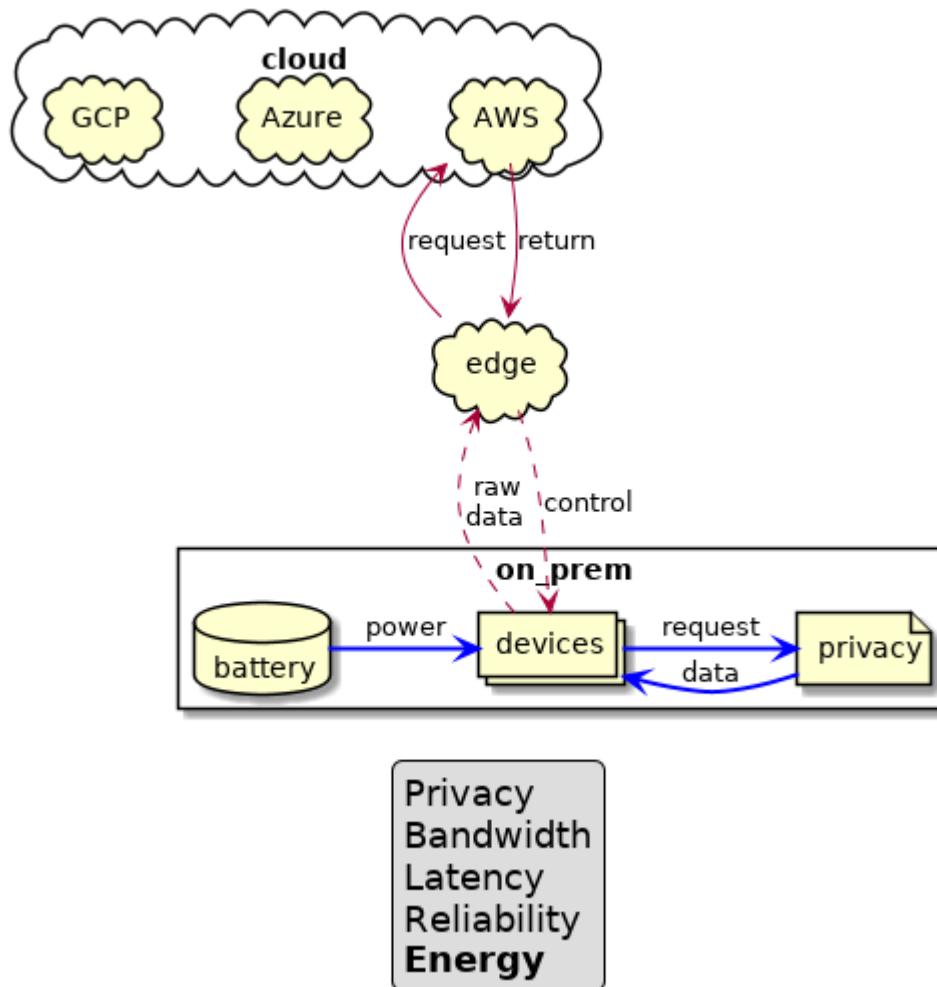
Scalability



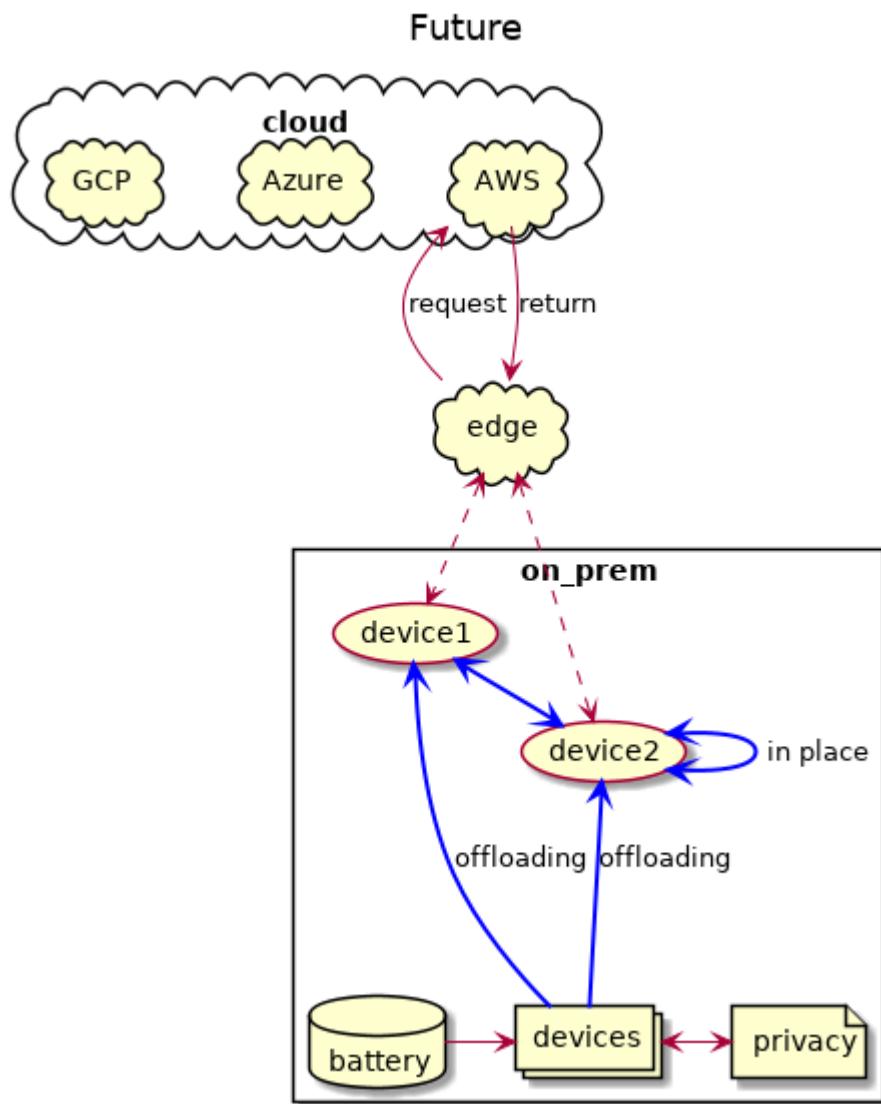
Currently



Edge computing problems



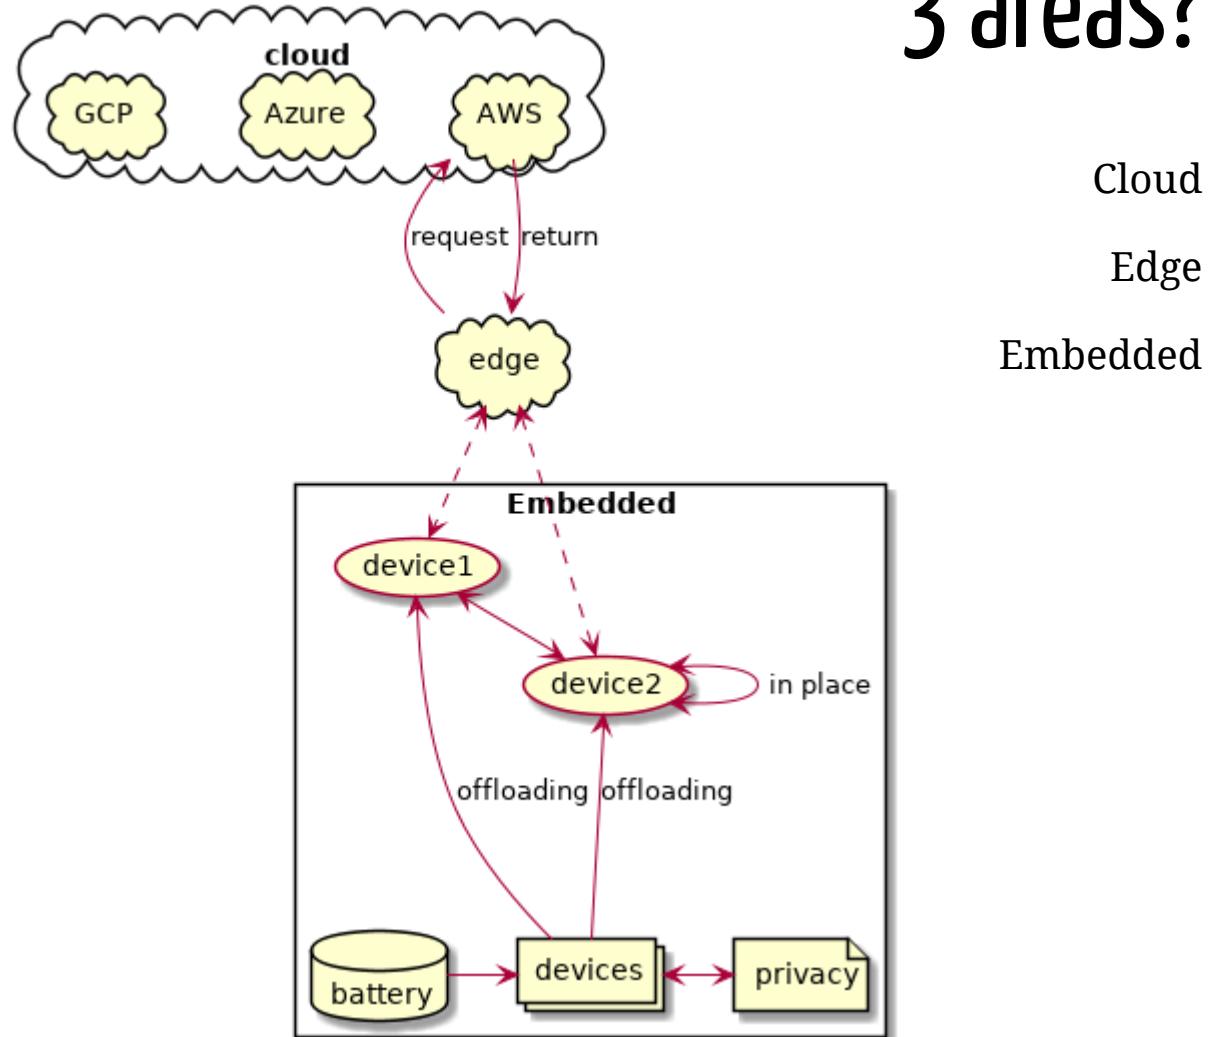
ref Tx is expensive.



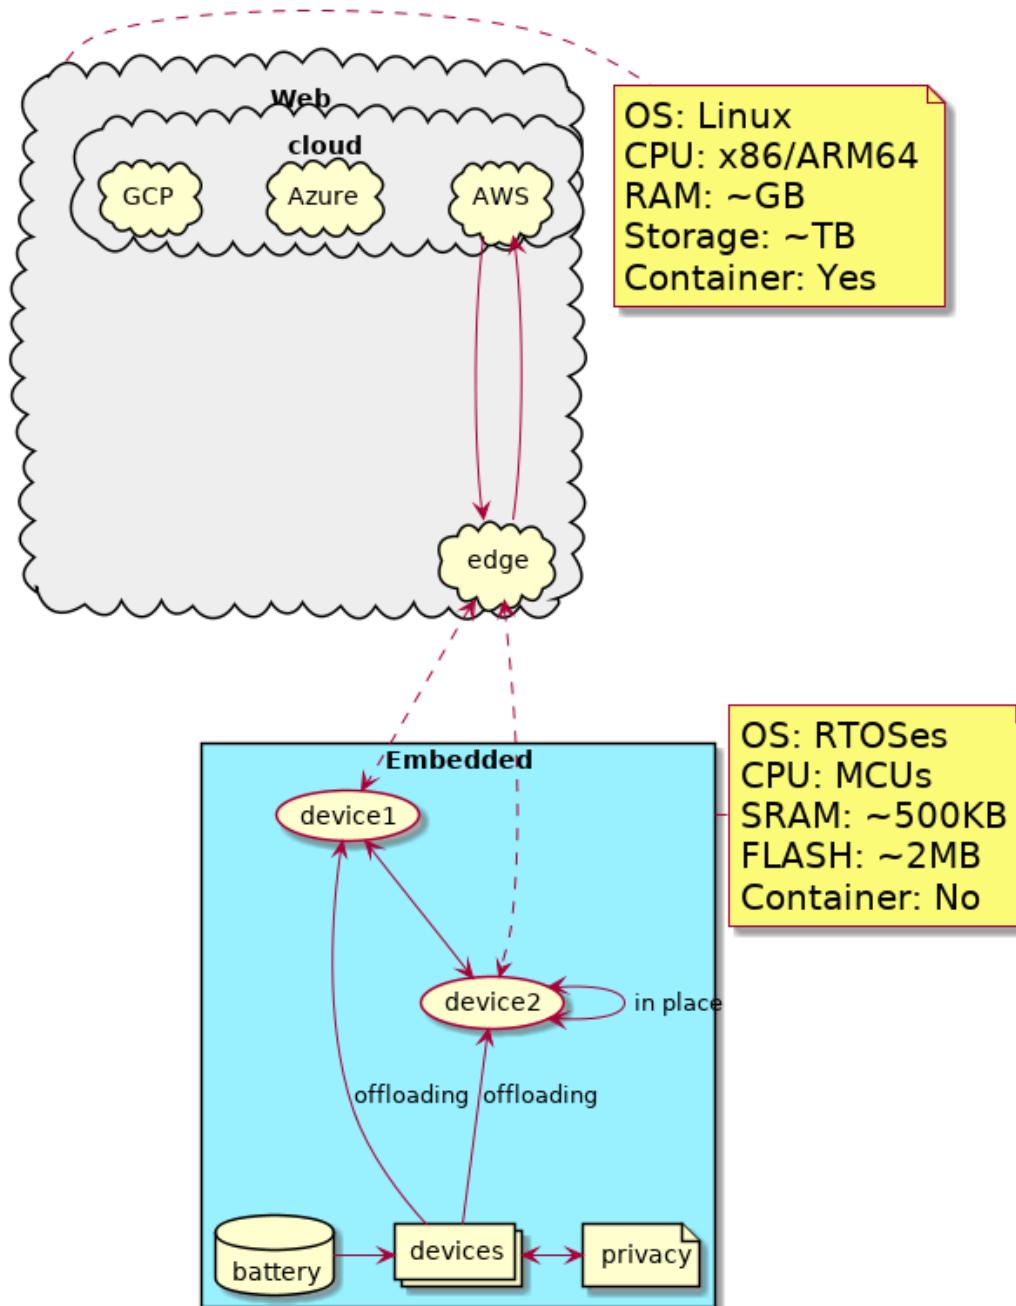
Outline

1. ~~Demo~~
2. **Problems**
 - ~~Edge Computing~~
 - **Web vs Embedded**
 - ML environment
3. Proposal
4. Three Enablers
5. PoC
6. Conclusion

3 areas?



Gap between Web and Embedded



Outline

1. ~~Demo~~

2. **Problems**

- ~~Edge Computing~~
- ~~Web vs Embedded~~
- **ML environment**

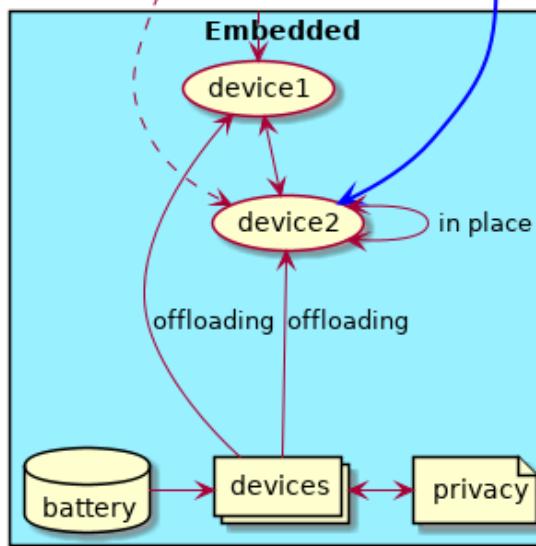
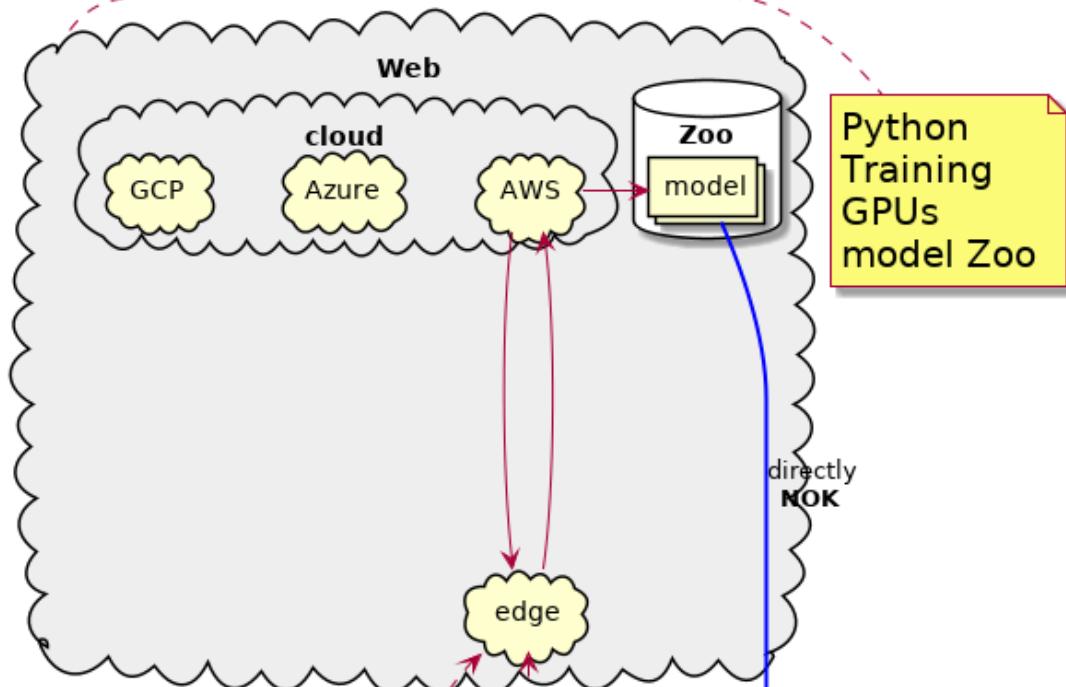
3. Proposal

4. Three Enablers

5. PoC

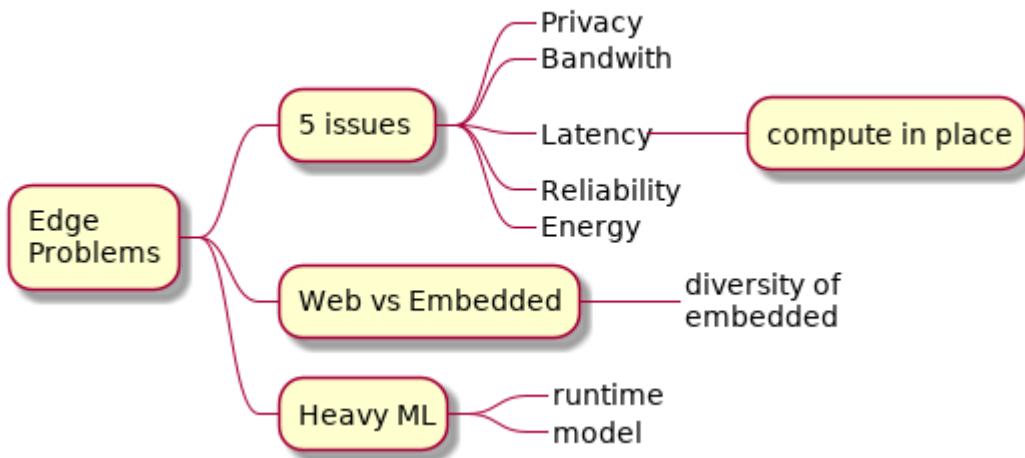
6. Conclusion

ML Problem



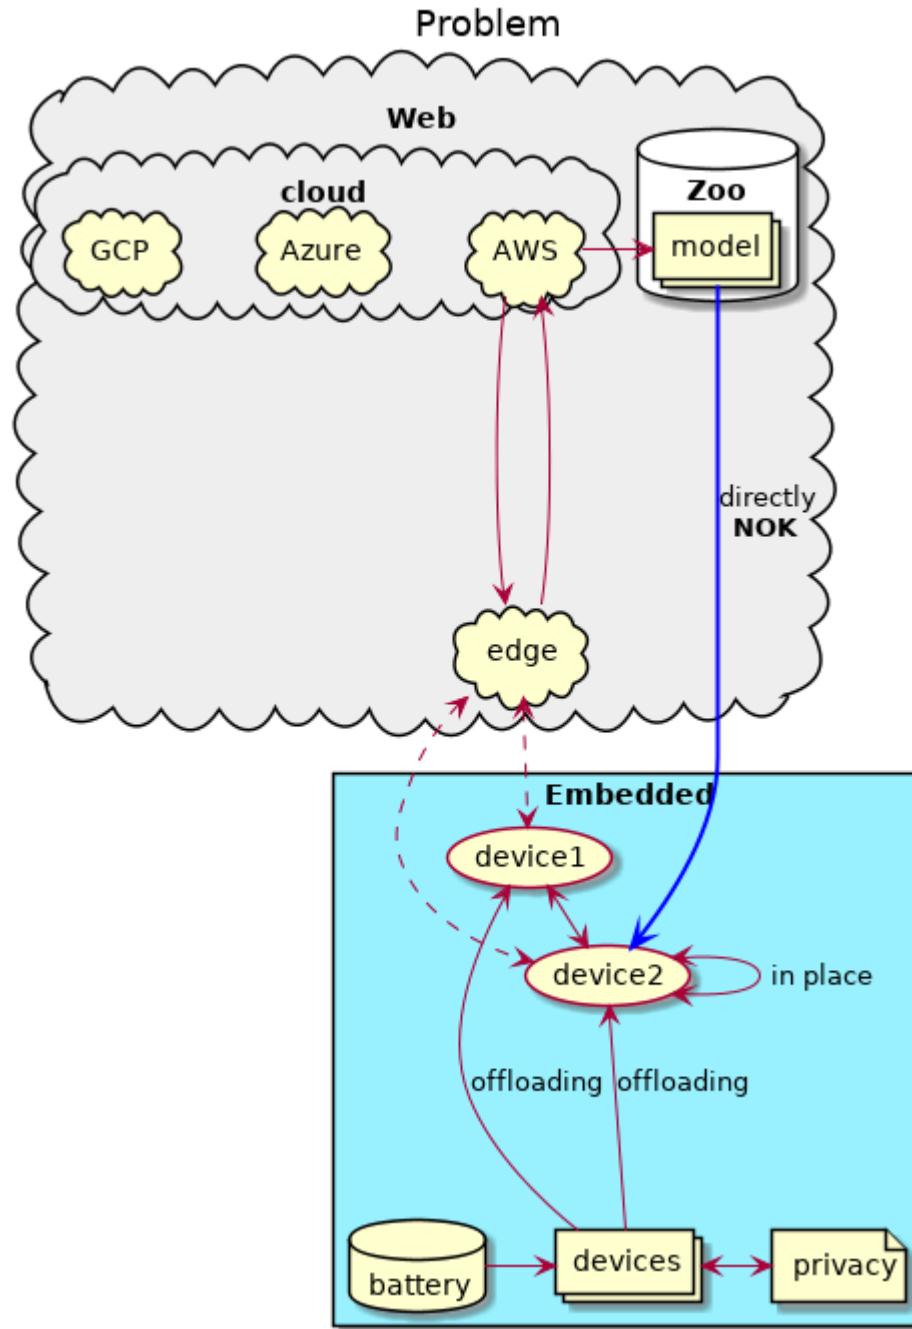
No Linux
No ML

Summary: Problems



Outline

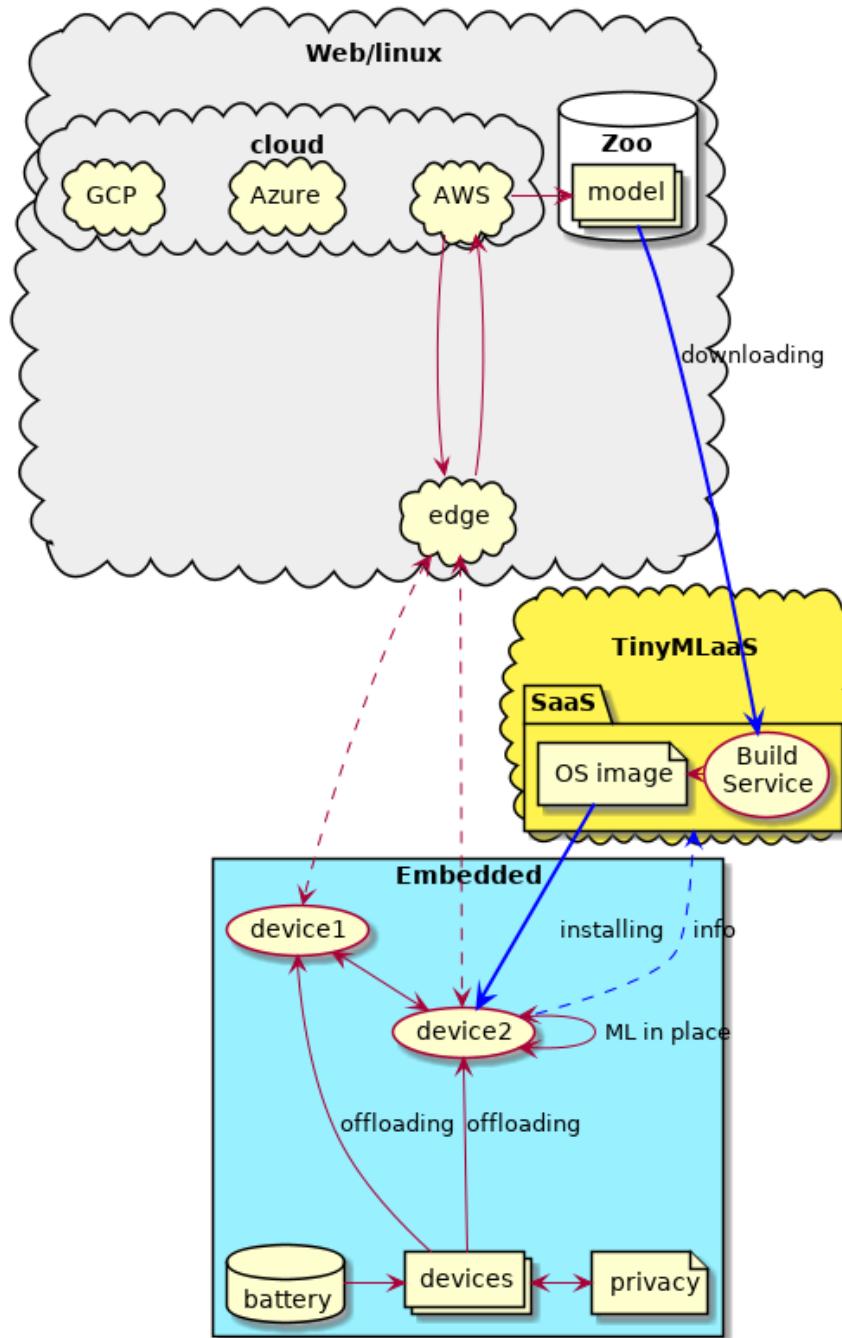
1. ~~Demo~~
2. ~~Problems~~
3. **Proposal**
4. Three Enablers
5. PoC
6. Conclusion

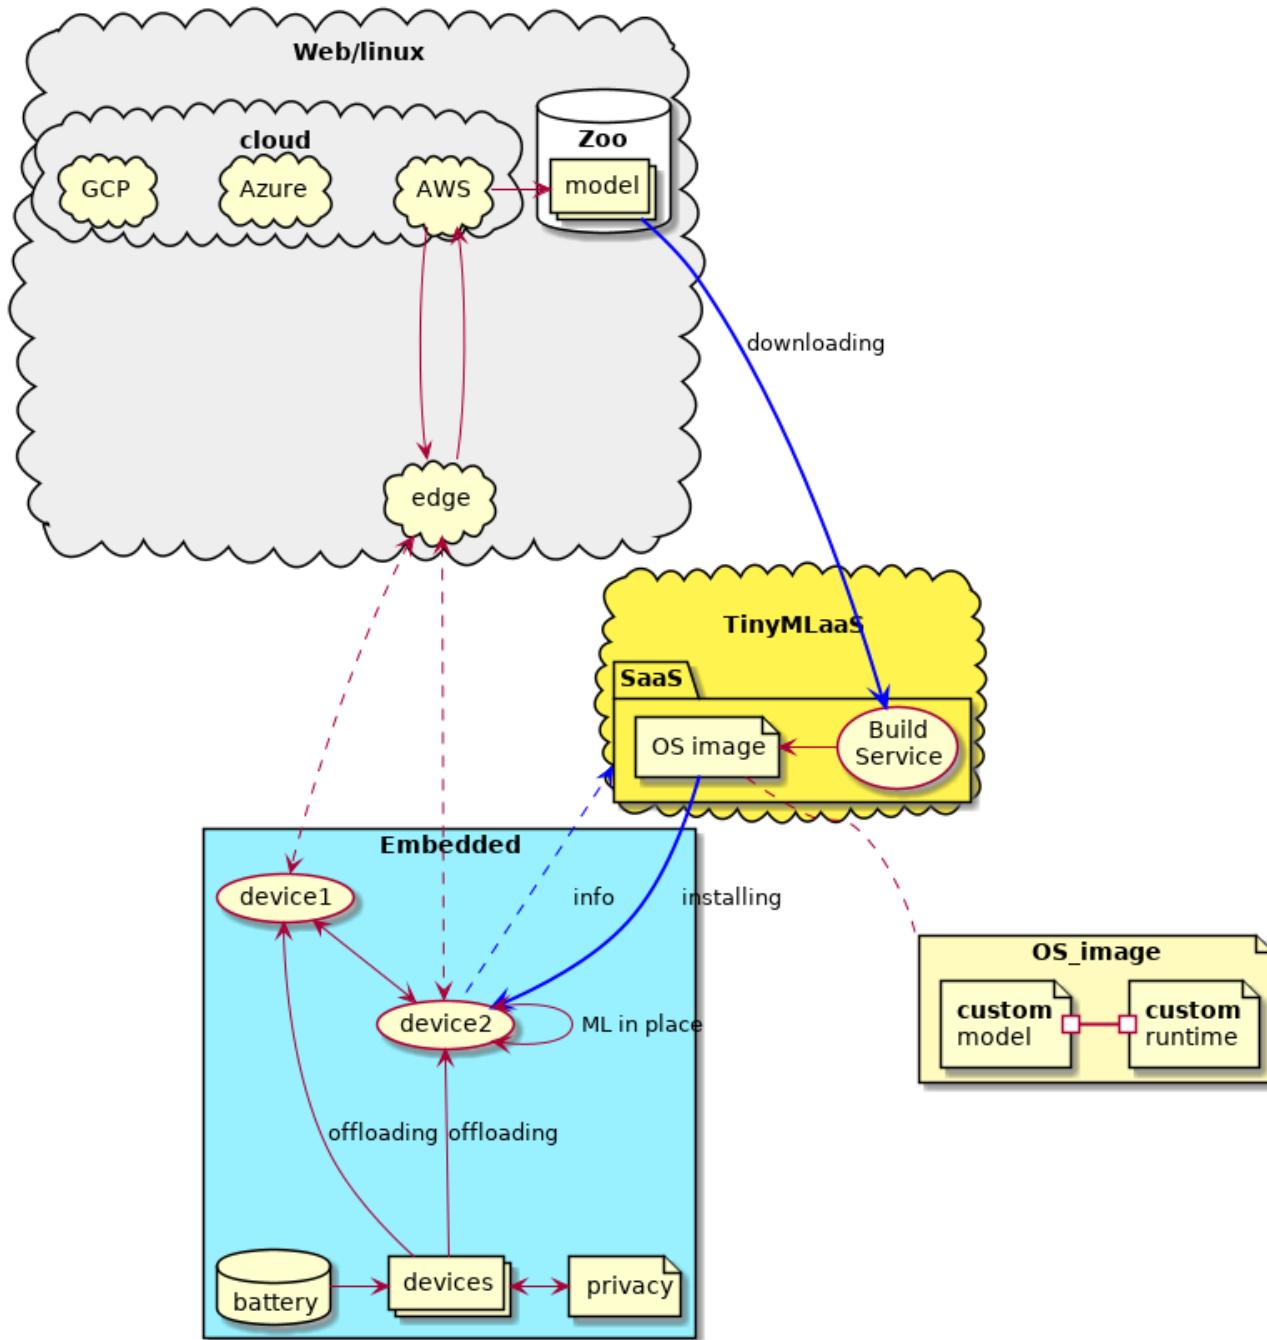


Squeeze ML

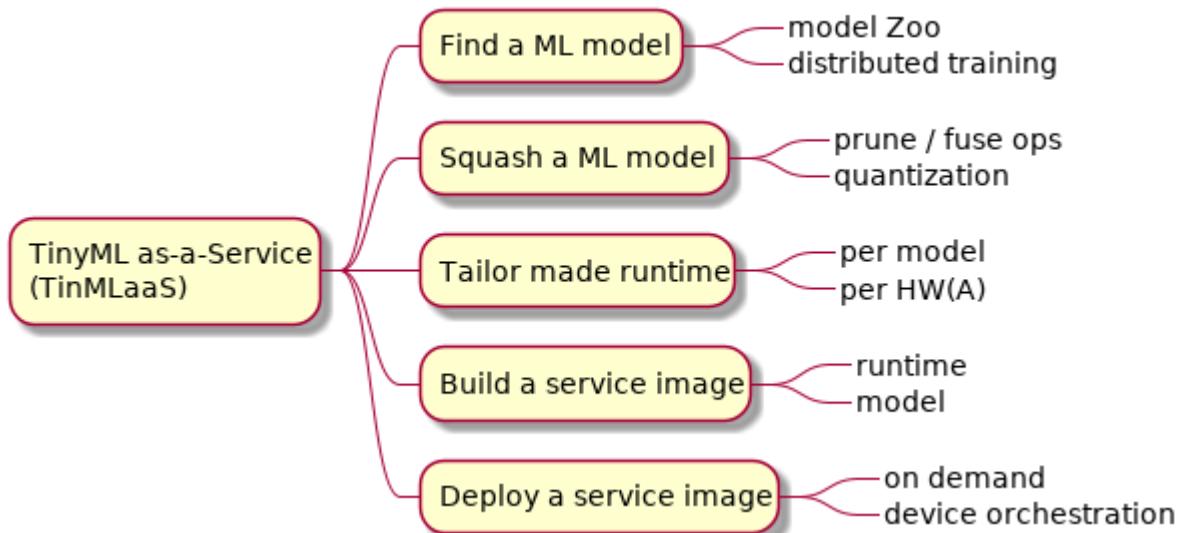
model & runtime

- per RTOS
- per MCU / ROM / RAM
- per HWA





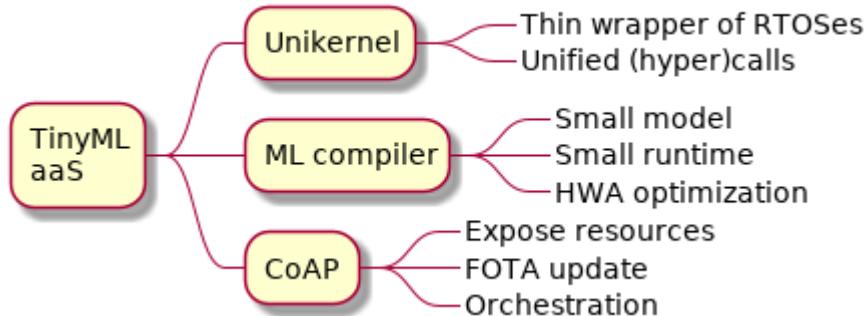
Proposal



Outline

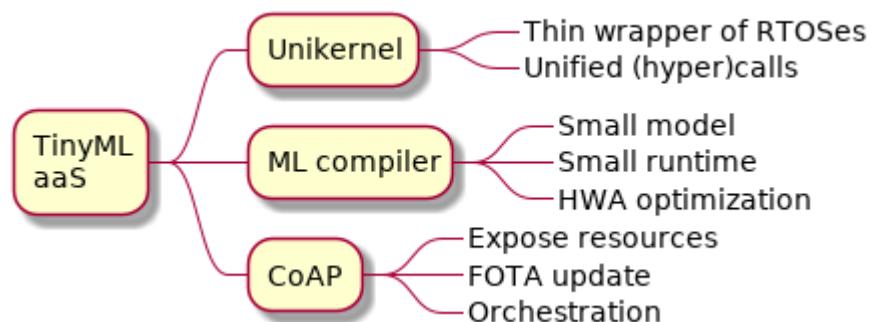
1. ~~Demo~~
2. ~~Problems~~
3. ~~Proposal~~
4. **Three Enablers**
5. PoC
6. Conclusion

Three enablers



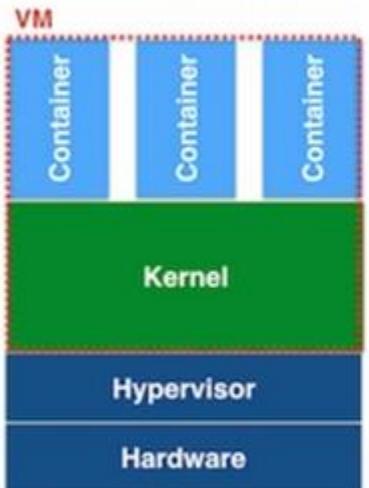
Outline

1. ~~Demo~~
2. ~~Problems~~
3. ~~Proposal~~
4. **Three Enablers**
 - **Unikernel**
 - ML compiler
 - CoAP
5. PoC
6. Conclusion

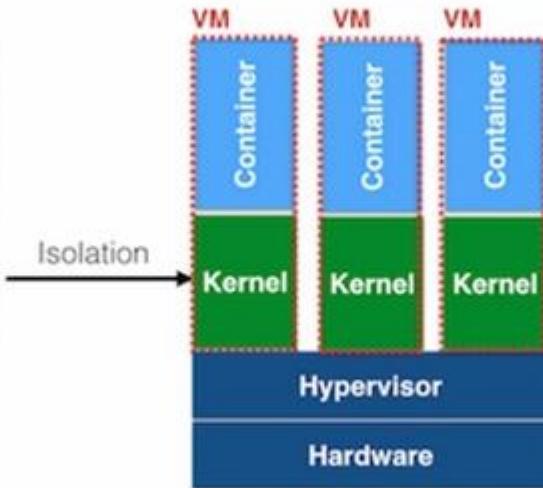


Unikernel

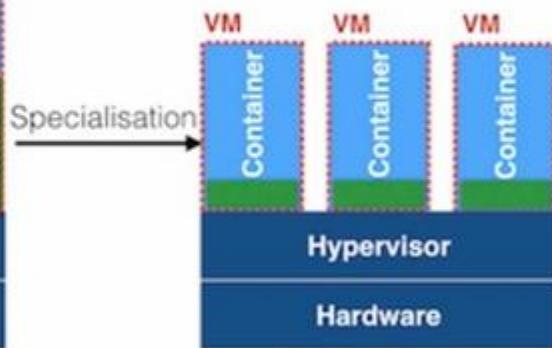
Isolation & specialisation with unikernels



Linux Container
Shared kernels



Container per VM
Duplicated kernels



Unikernels
Specialized kernels

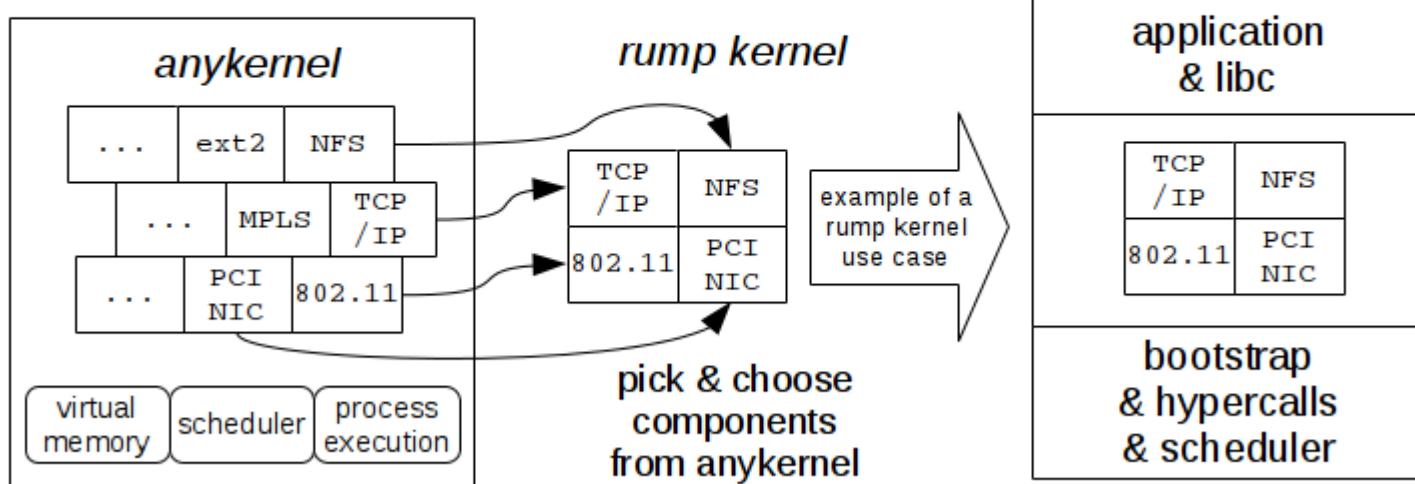
Isolation

Specialisation

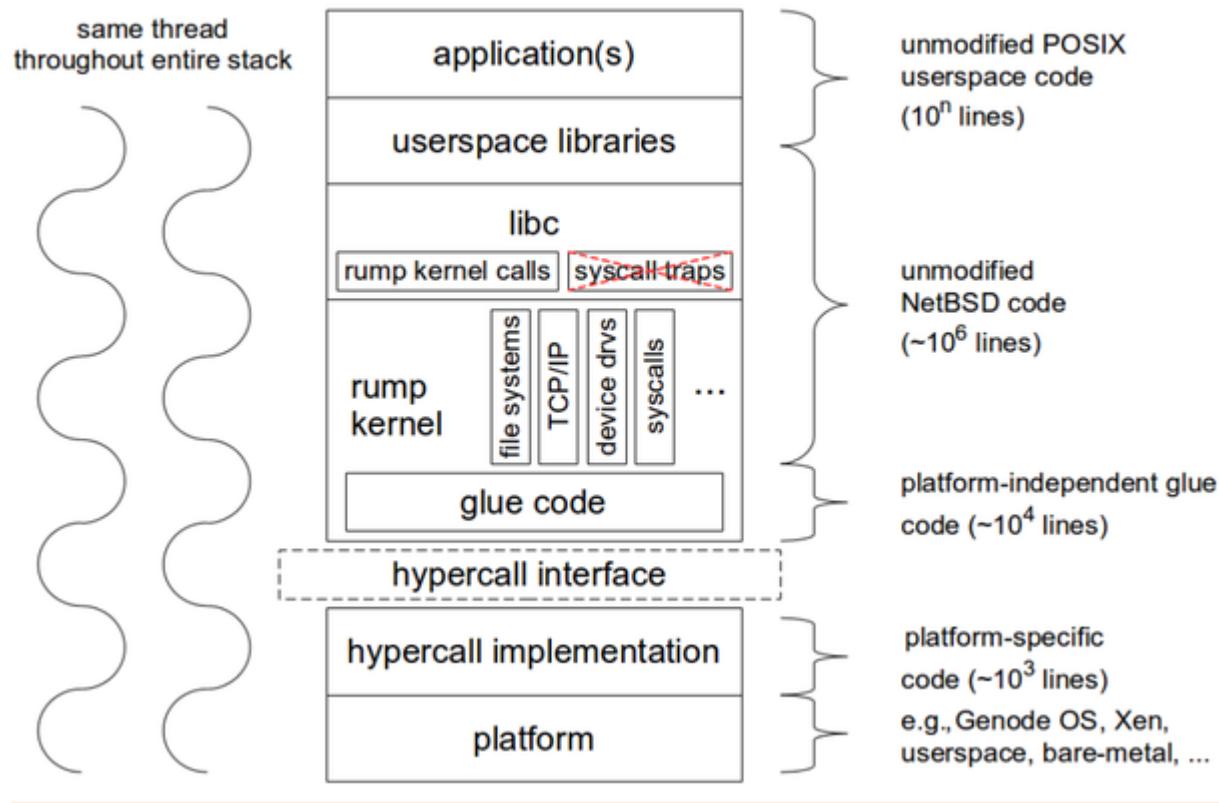
library Operating System (libOS)

The relationship between *anykernel*,
rump kernel and *unikernel*

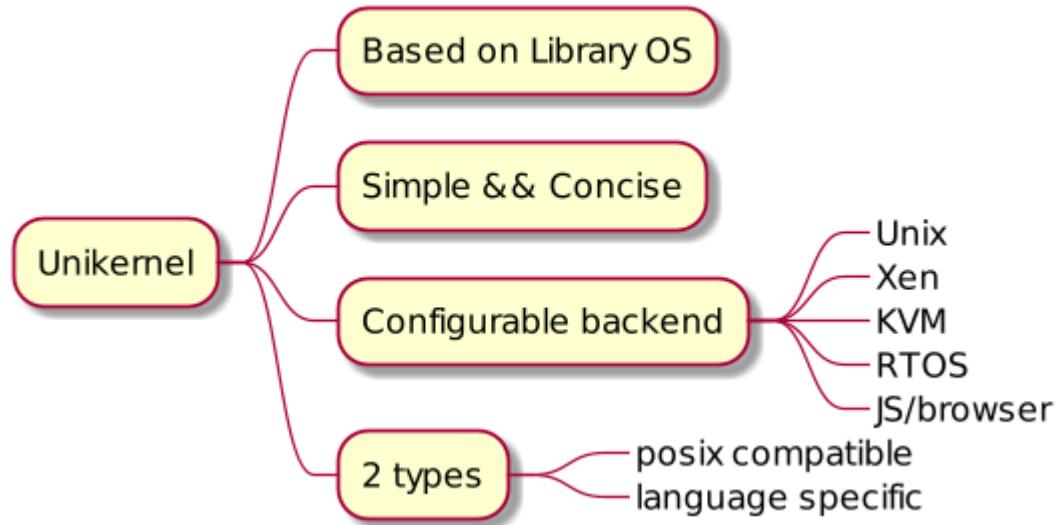
Rumprun unikernel



Internal



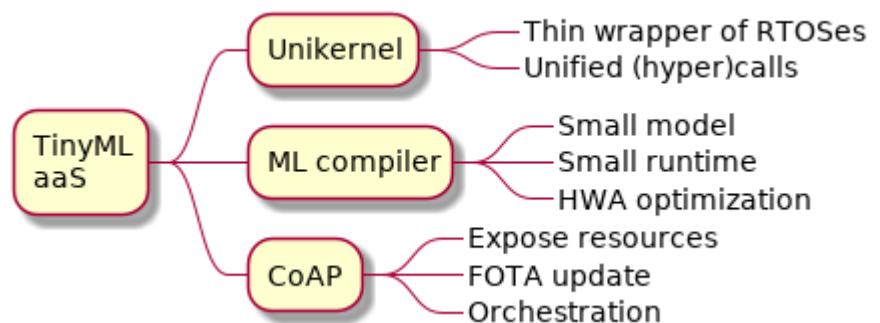
Summary: Unikernel



Can be a thin **wrapper** of different RTOSes?

Outline

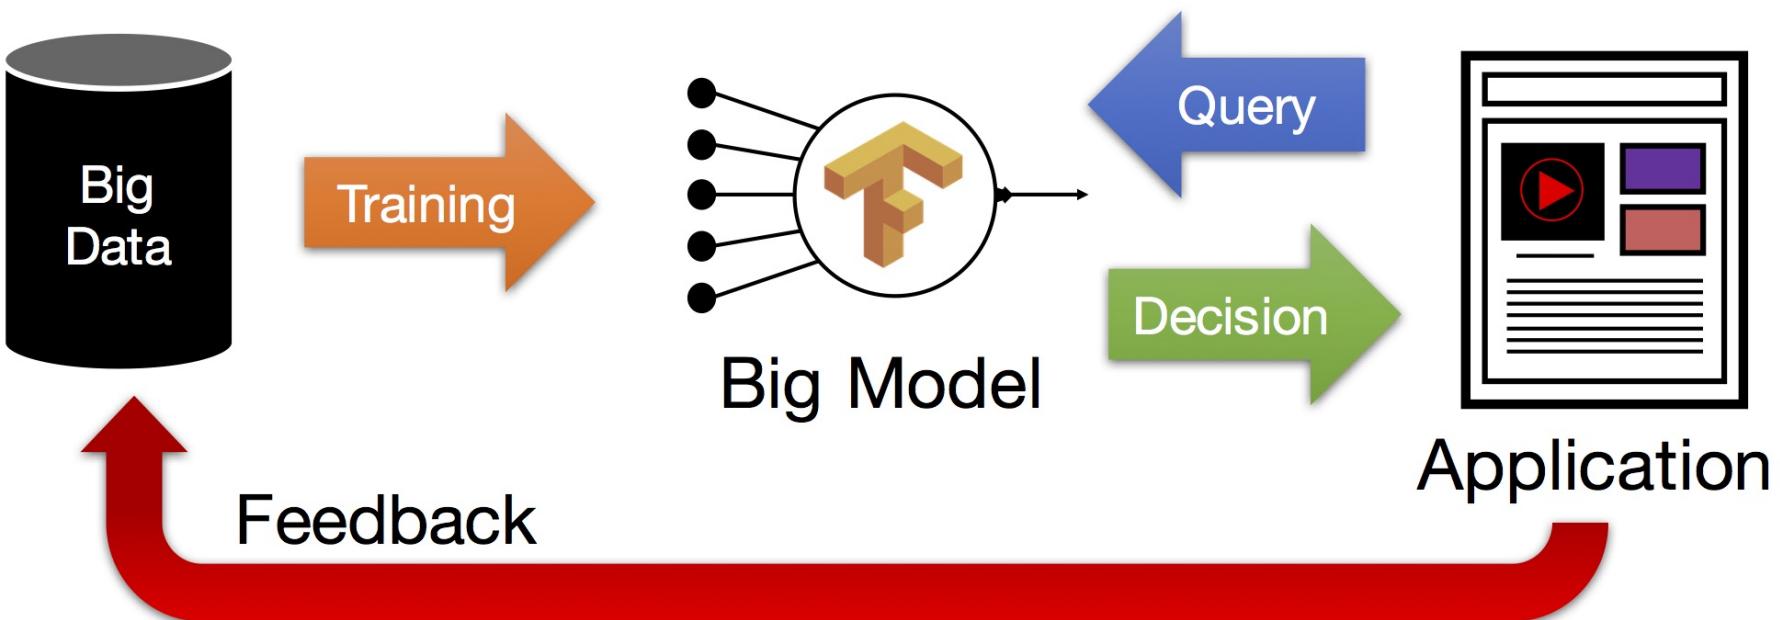
1. ~~Demo~~
2. ~~Problems~~
3. ~~Proposal~~
4. **Three Enablers**
 - **Unikernel**
 - **ML compiler**
 - CoAP
5. PoC
6. Conclusion



2 ML phases:

Learning

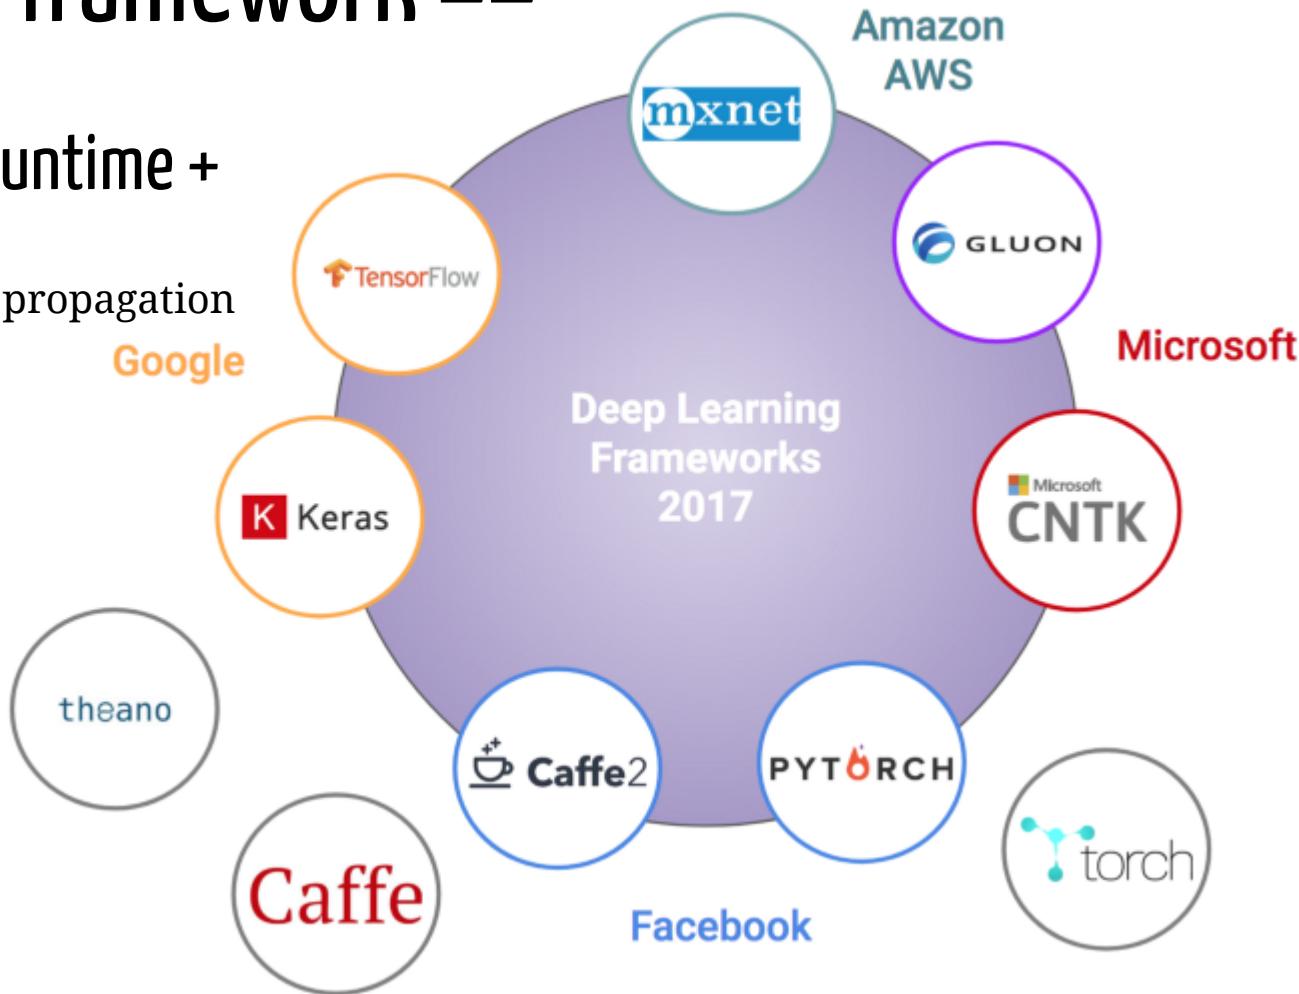
Inference



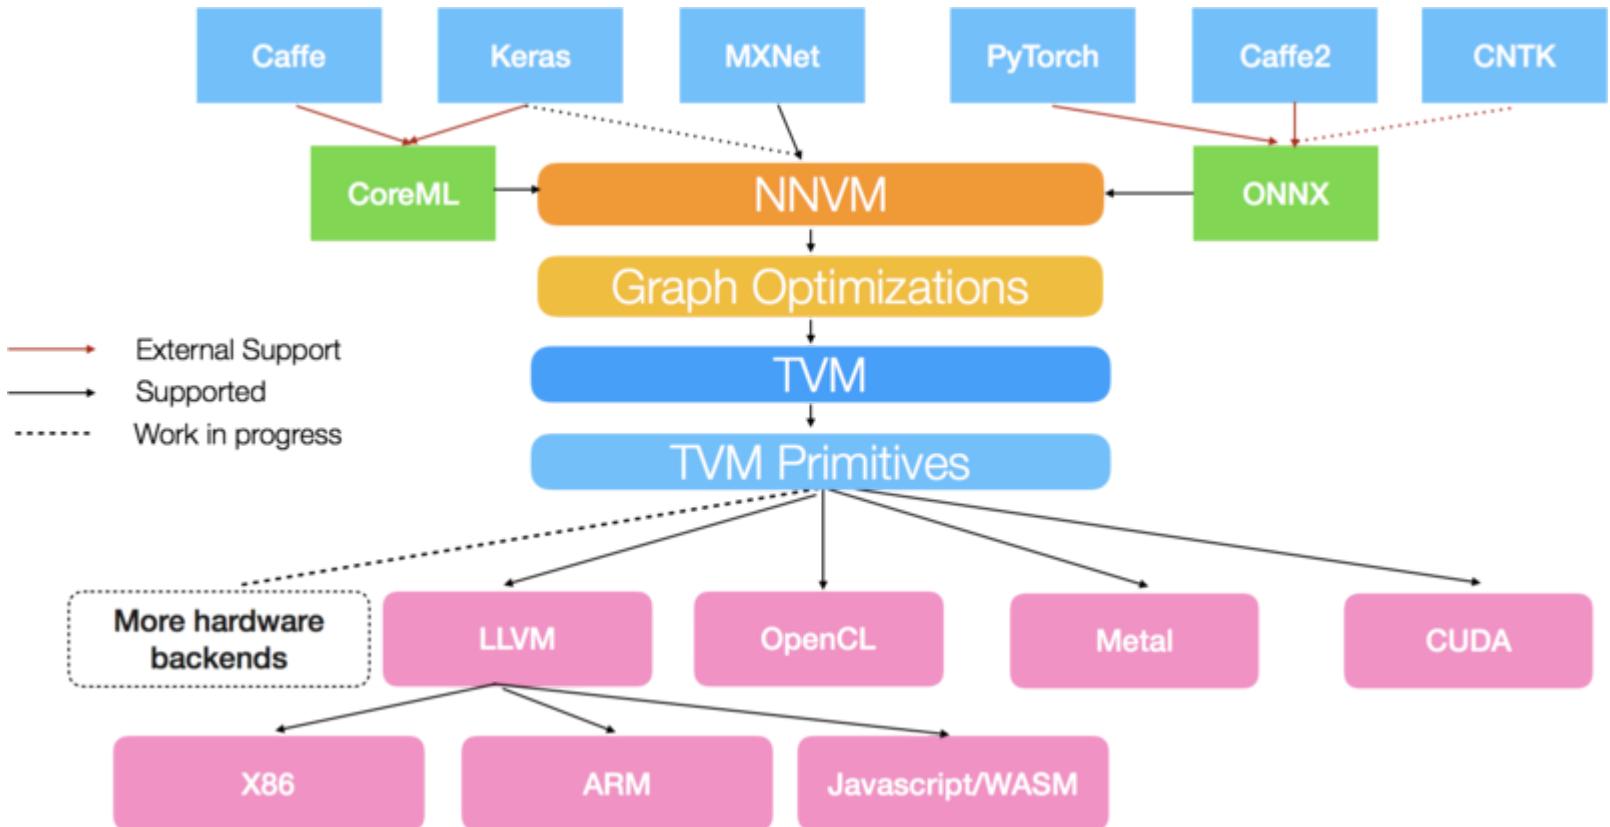
ML framework ==

ML runtime +

back propagation

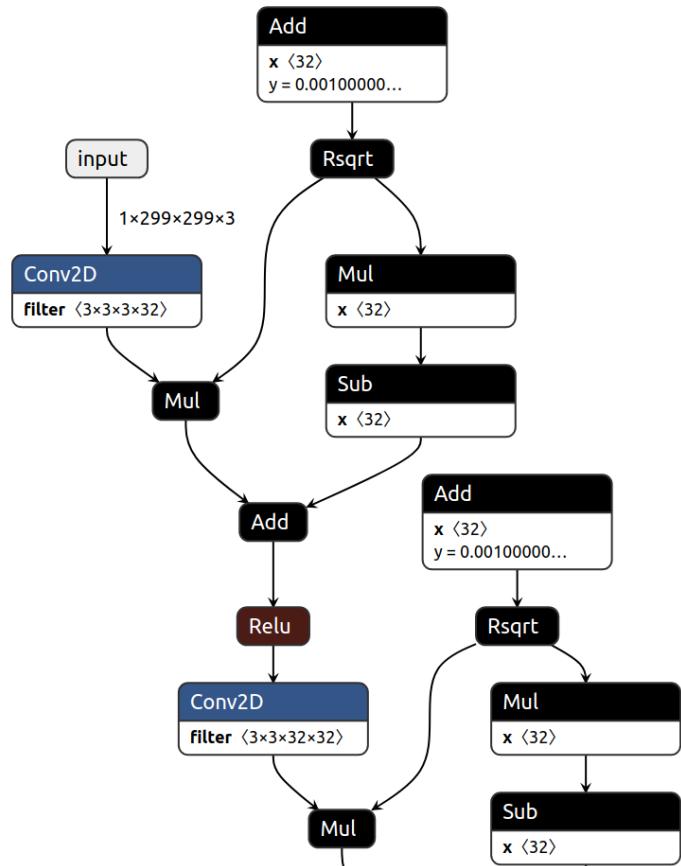


Compiling



Computational graph (cgraph)

[≡] [+]



NODE PROPERTIES

type Conv2D

name InceptionV3/InceptionV3/Conv2d_2a_3x3/convolution

X

ATTRIBUTES

data_format NHWC

padding VALID

strides 1, 1, 1, 1

T Float32

use_cudnn_on_g... true

INPUTS

input id: InceptionV3/InceptionV3/Conv2d_1a_3x3/Relu

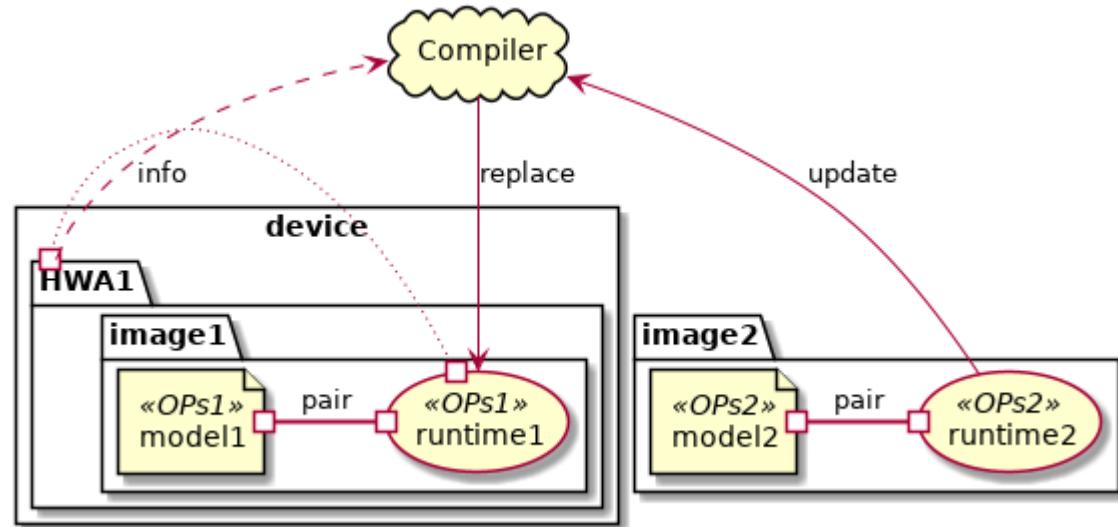
filter id: InceptionV3/Conv2d_2a_3x3/weights/read

OUTPUTS

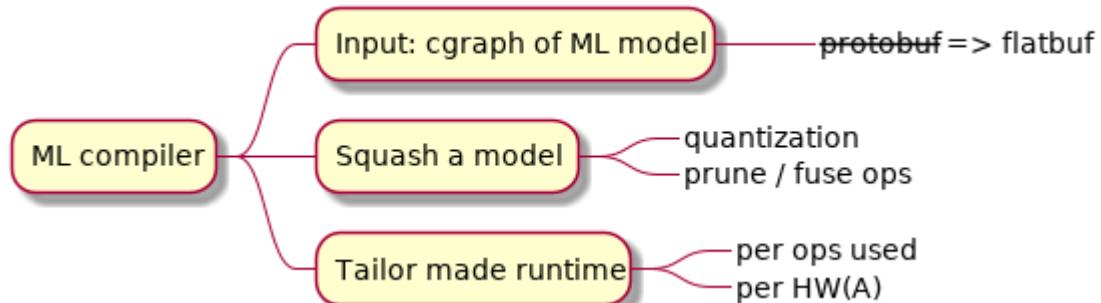
output id: InceptionV3/InceptionV3/Conv2d_2a_3x3/convolution

General vs Special purpose runtime

- General purpose is too big
 - all OPs built-in
- Model could be optimized,
 - per HWAs' OPs,
 - smaller than CPU's OPs(?)
- Runtime should implement **only** OPs,
 - which model uses.

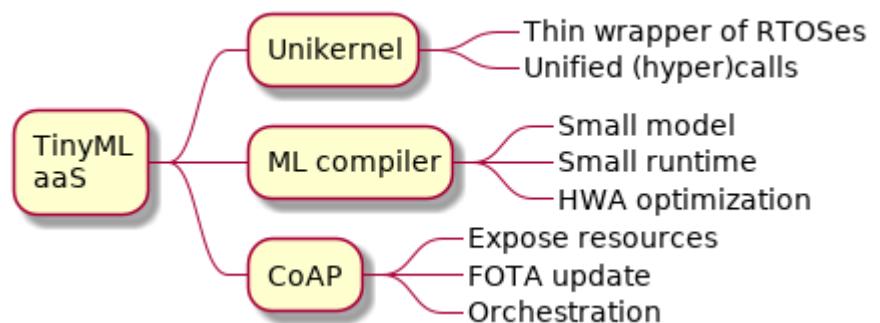


Summary: ML compiler



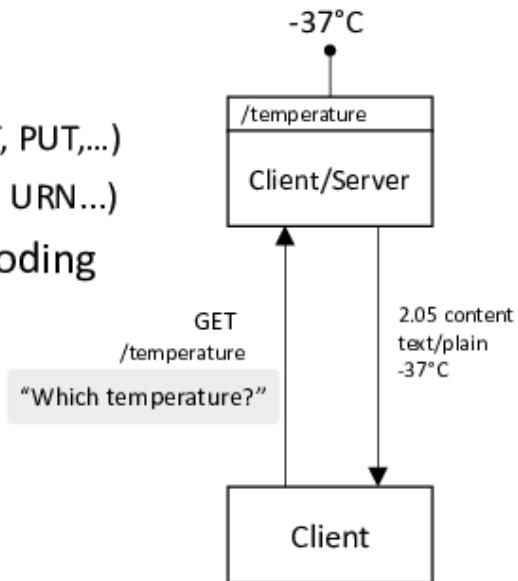
Outline

1. ~~Demo~~
2. ~~Problems~~
3. ~~Proposal~~
4. **Three Enablers**
 - o ~~Unikernel~~
 - o ~~ML compiler~~
 - o ~~CoAP~~
5. PoC
6. Conclusion

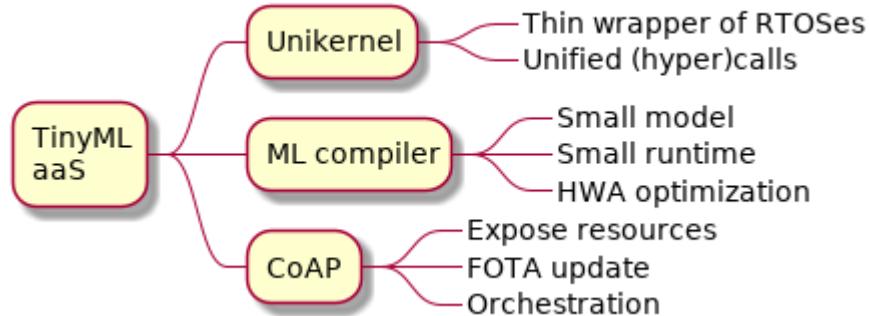


The Constrained Application Protocol (CoAP)

- CoAP (RFC7252) implements HTTP's **REST** model
 - Simple devices: 100 to 250 KiB code and 10 to 50 KiB RAM
 - Each device can be client and server exposing resources
 - CoAP defines methods to access those resources (GET, POST, PUT,...)
 - Same key concepts borrowed from HTTP (Media types, URL, URN...)
- Has a compact 4-byte header, with simple options encoding
- Simple protocol, datagram (UDP, DTLS)
 - Reliability through header message type "*CON/NON*"
 - With TCP/TLS (RFC8323) support for NAT-ed environments
- The Resource Directory provides a directory service



Summary: Three Enablers



Outline

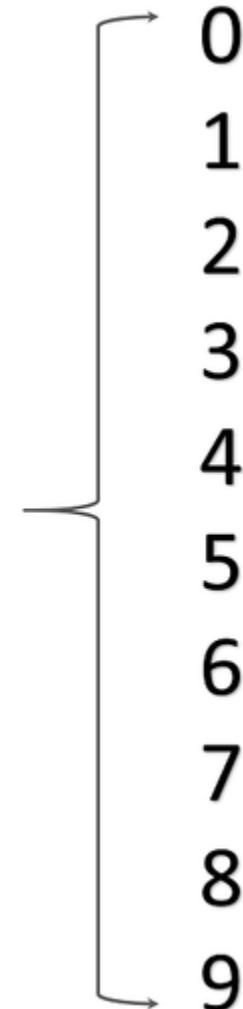
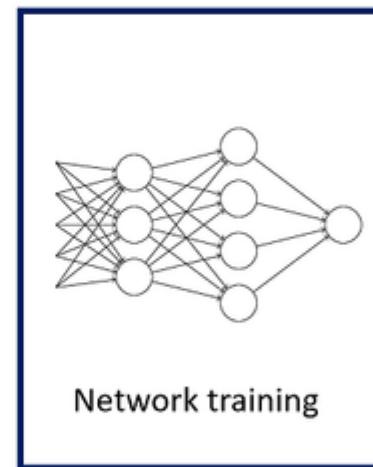
1. ~~Demo~~
2. ~~Problems~~
3. ~~Proposal~~
4. ~~Three Enablers~~
5. PoC
6. Conclusion

MNIST: Handwriting digits recognition

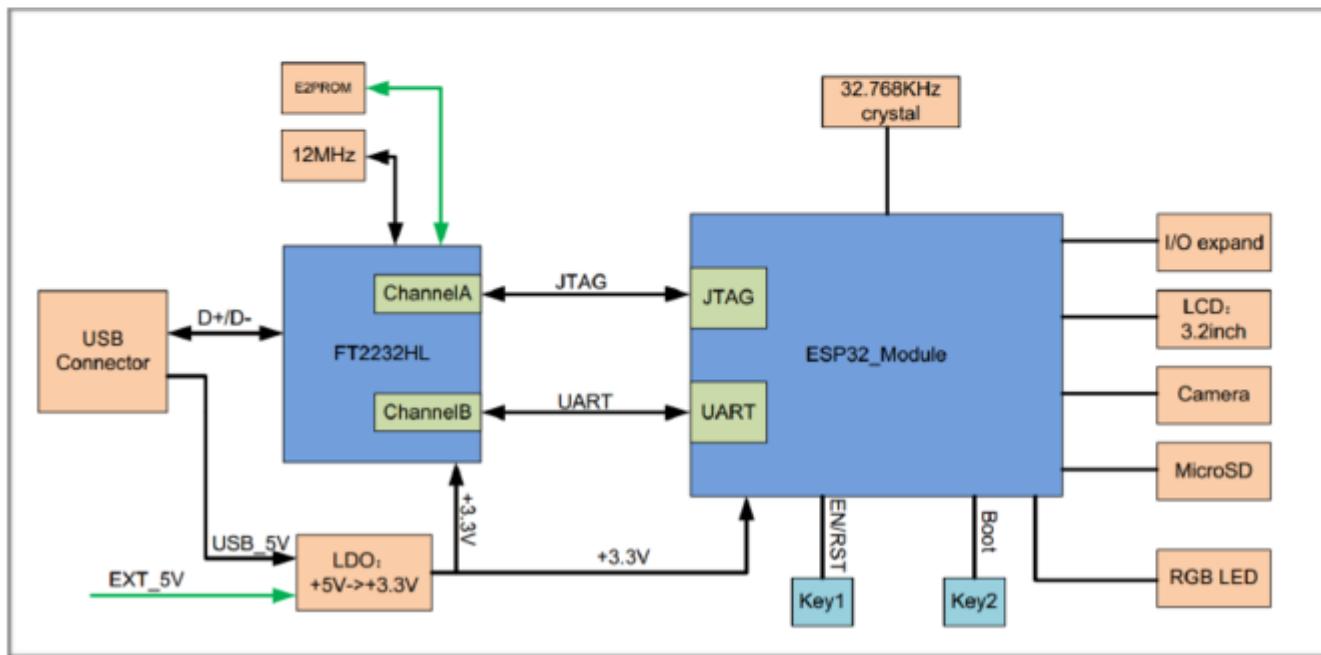
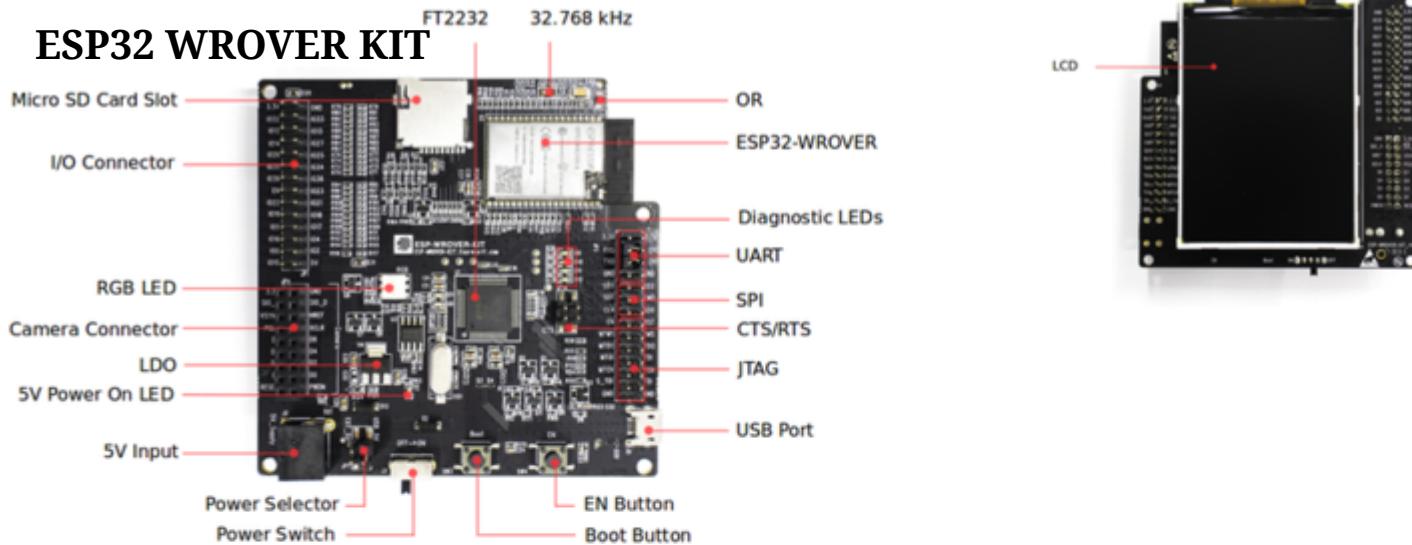
60K images for training, 10K for testing

0 0 0 0 0 0 0 0 0 0 0 0 0
1 1 1 1 1 1 1 1 1 1 1 1 1
2 2 2 2 2 2 2 2 2 2 2 2 2
3 3 3 3 3 3 3 3 3 3 3 3 3
4 4 4 4 4 4 4 4 4 4 4 4 4
5 5 5 5 5 5 5 5 5 5 5 5 5
6 6 6 6 6 6 6 6 6 6 6 6 6
7 7 7 7 7 7 7 7 7 7 7 7 7
8 8 8 8 8 8 8 8 8 8 8 8 8
9 9 9 9 9 9 9 9 9 9 9 9 9

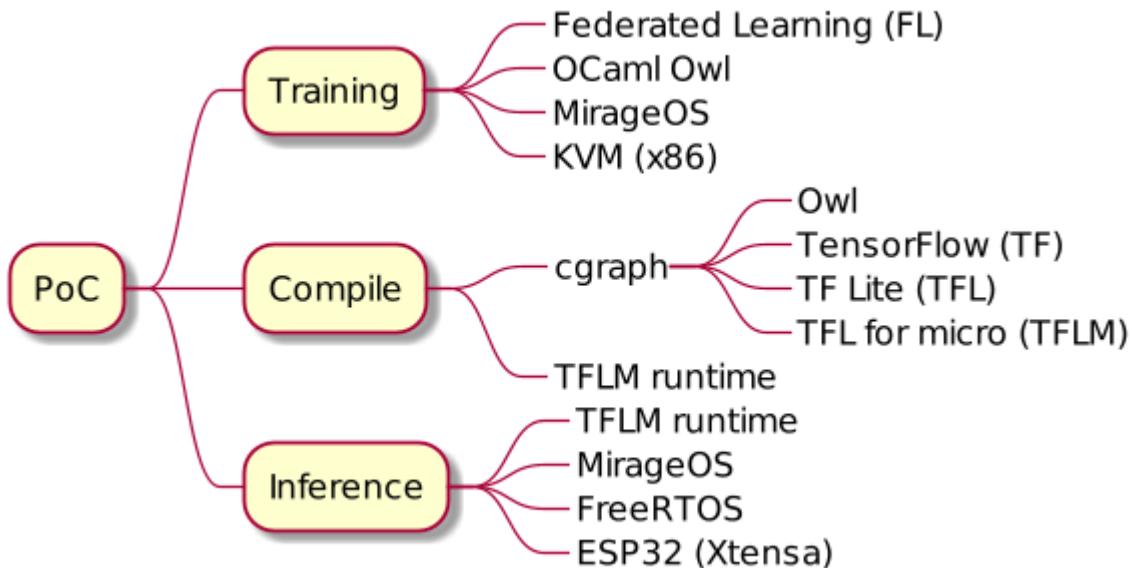
Data & Labels



ESP32 WROVER KIT

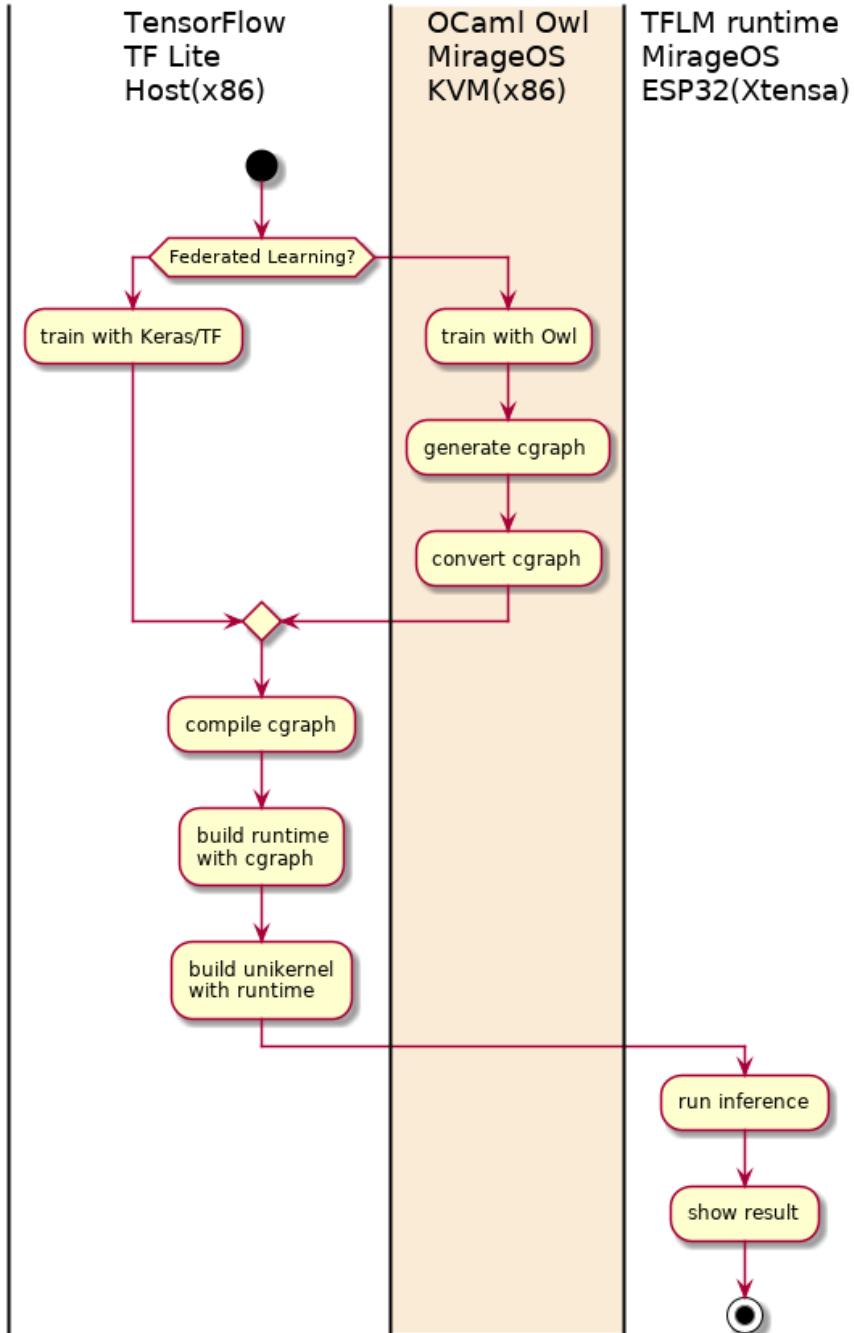


PoC Outline



PoC

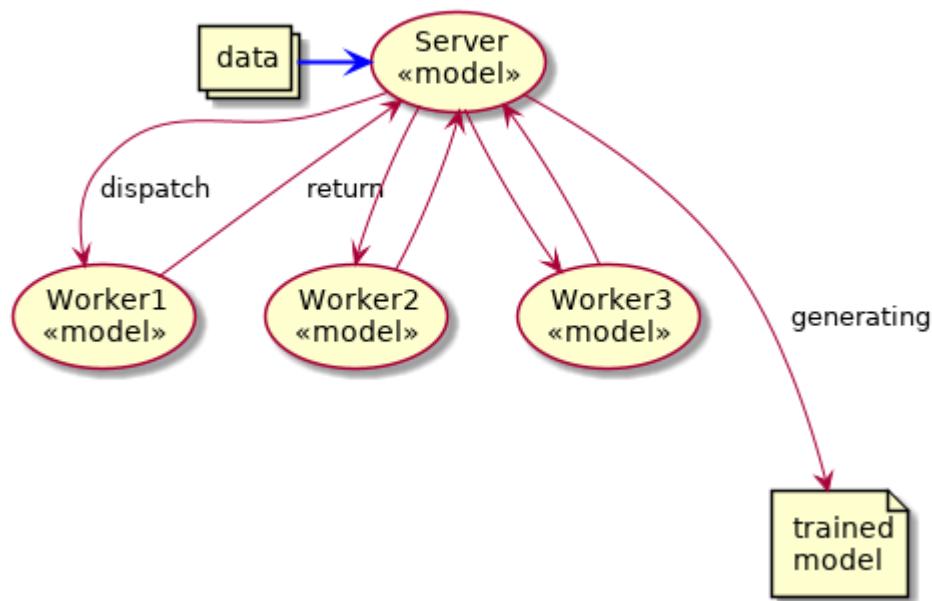
Sequence



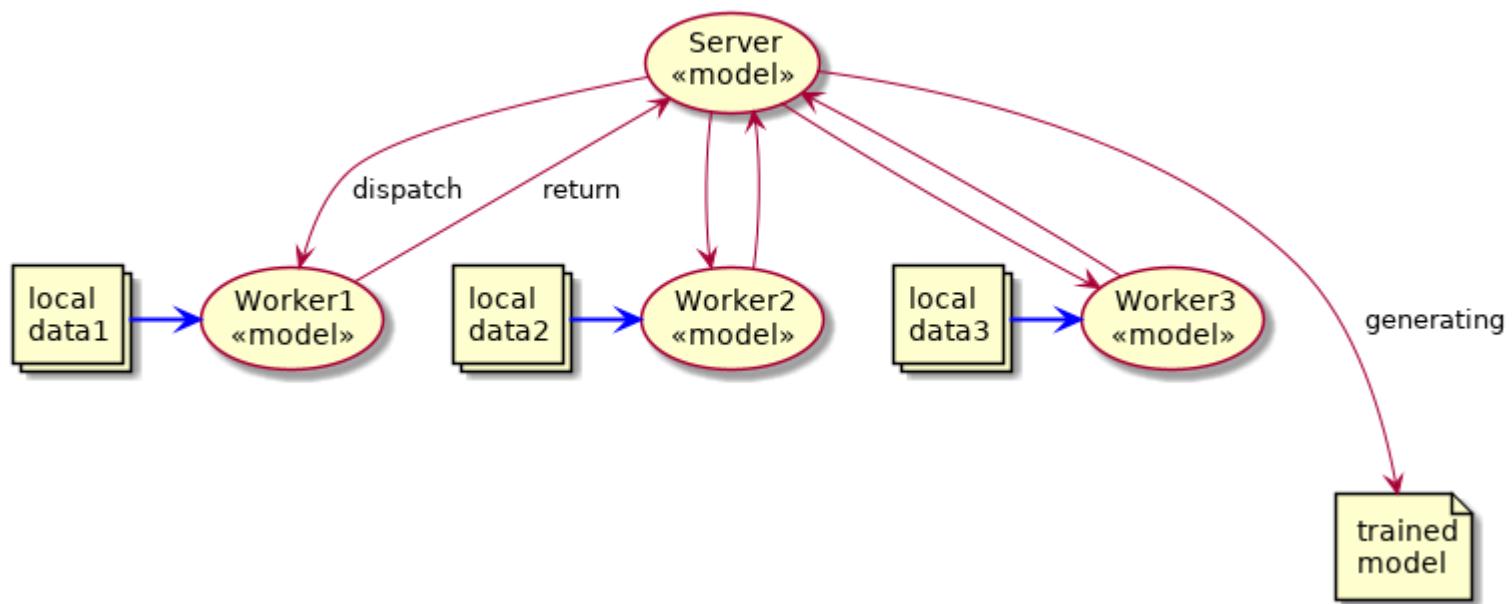
Outline

1. ~~Demo~~
2. ~~Problems~~
3. ~~Proposal~~
4. ~~Three Enablers~~
5. PoC
 - Training
 - Compile
 - Inference
6. Conclusion

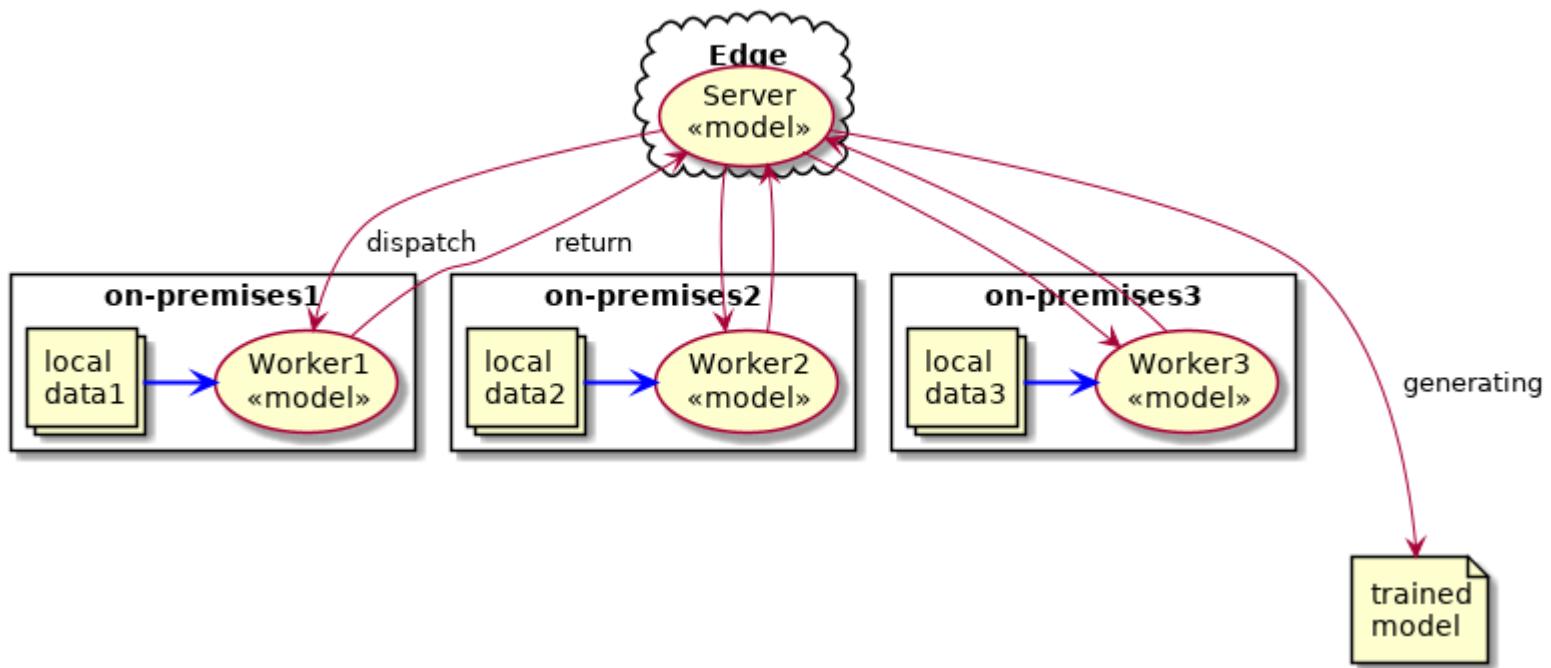
Training: Parameter Server (PS)

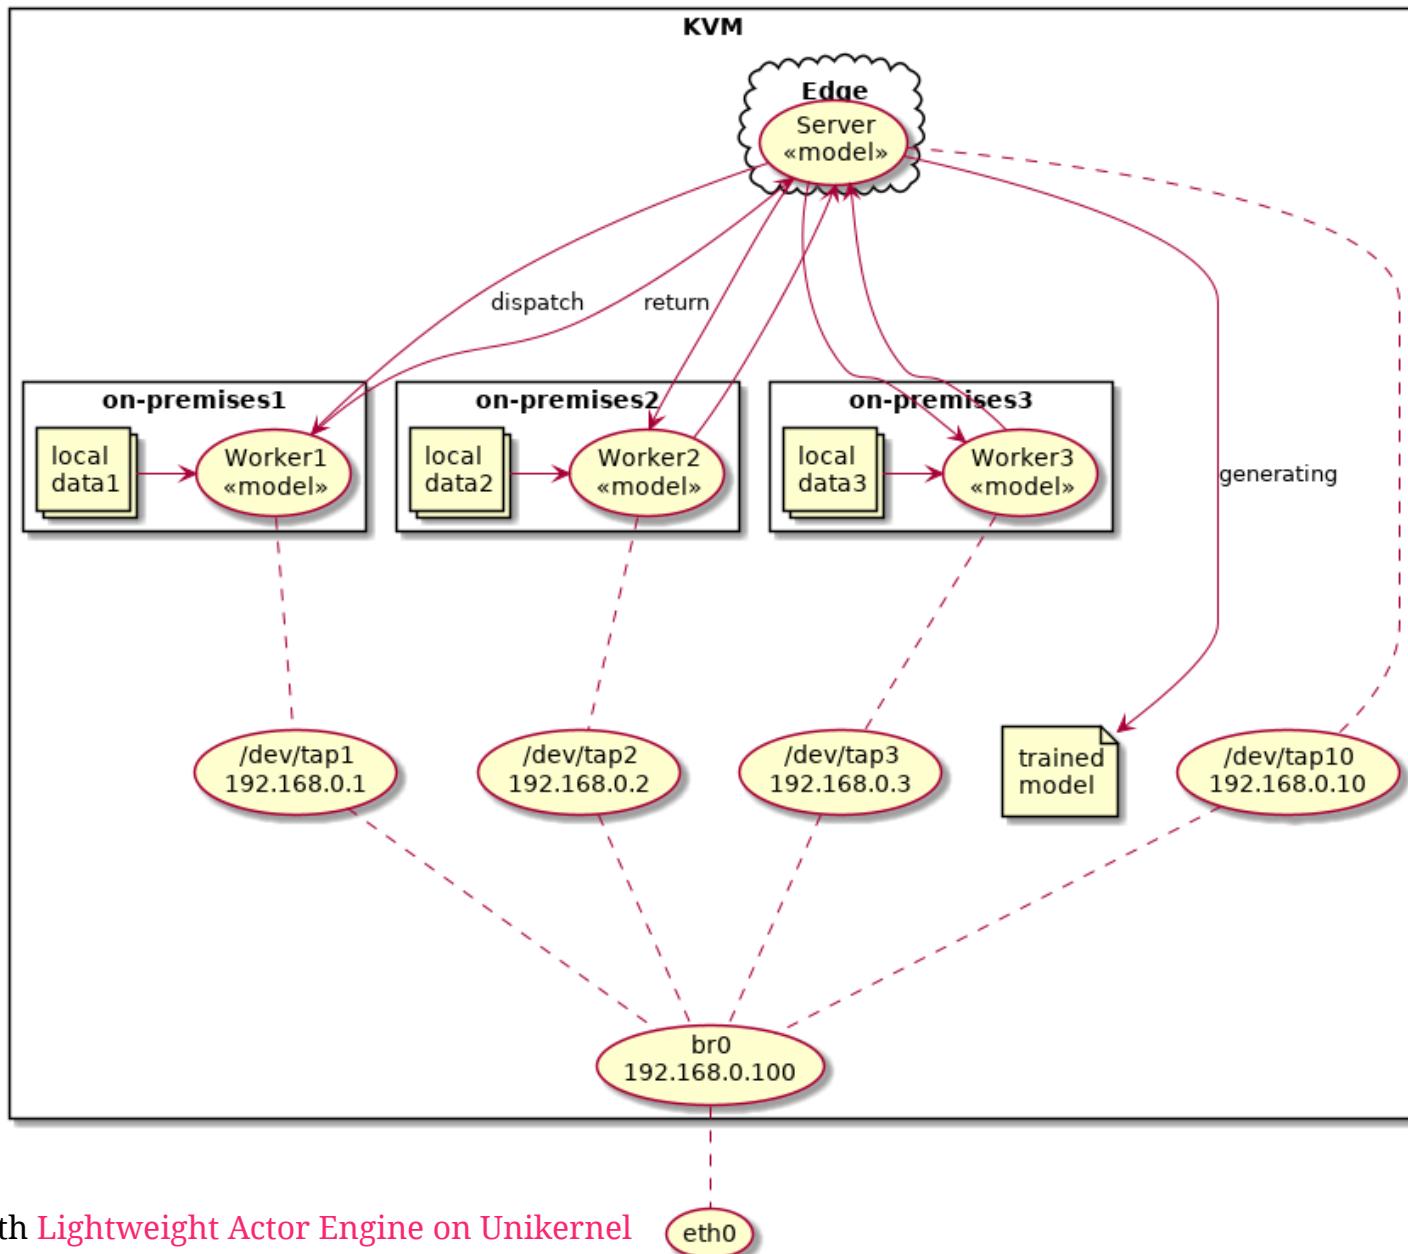


Training: Federated Learning (FL)



Training: Federated Learning (FL)

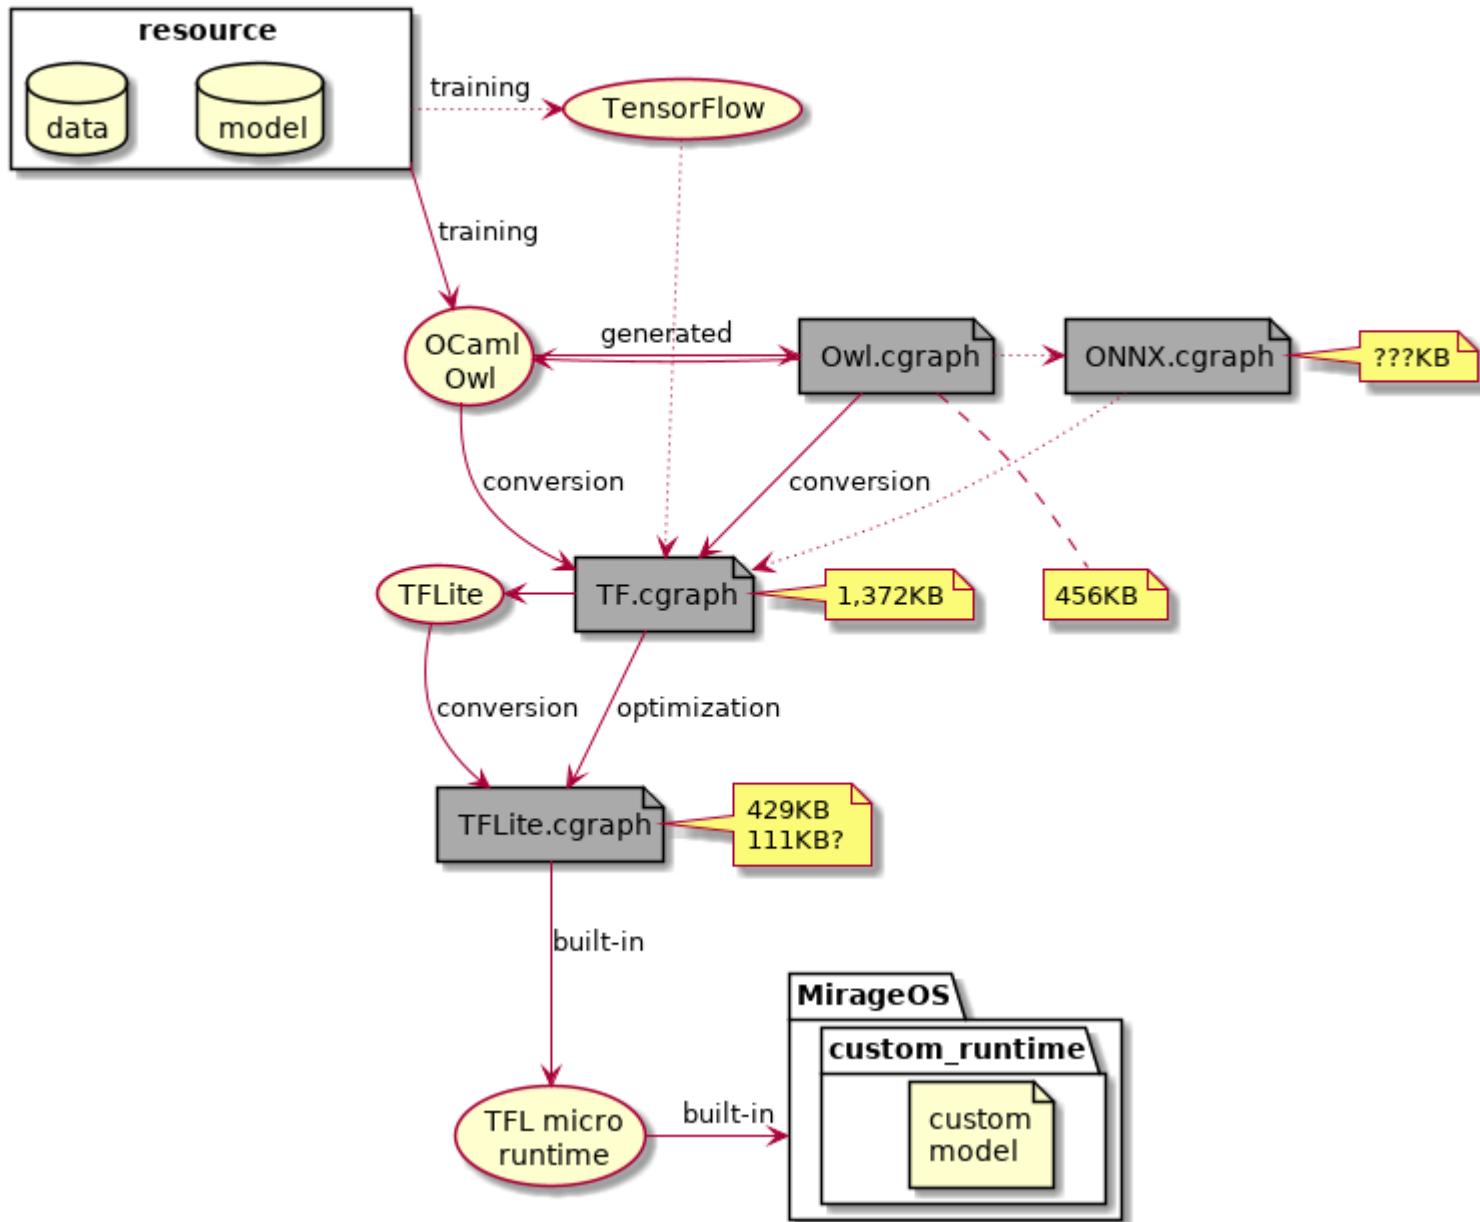




Outline

1. ~~Demo~~
2. ~~Problems~~
3. ~~Proposal~~
4. ~~Three Enablers~~
5. PoC
 - ~~Training~~
 - ~~Compile~~
 - Inference
6. Conclusion

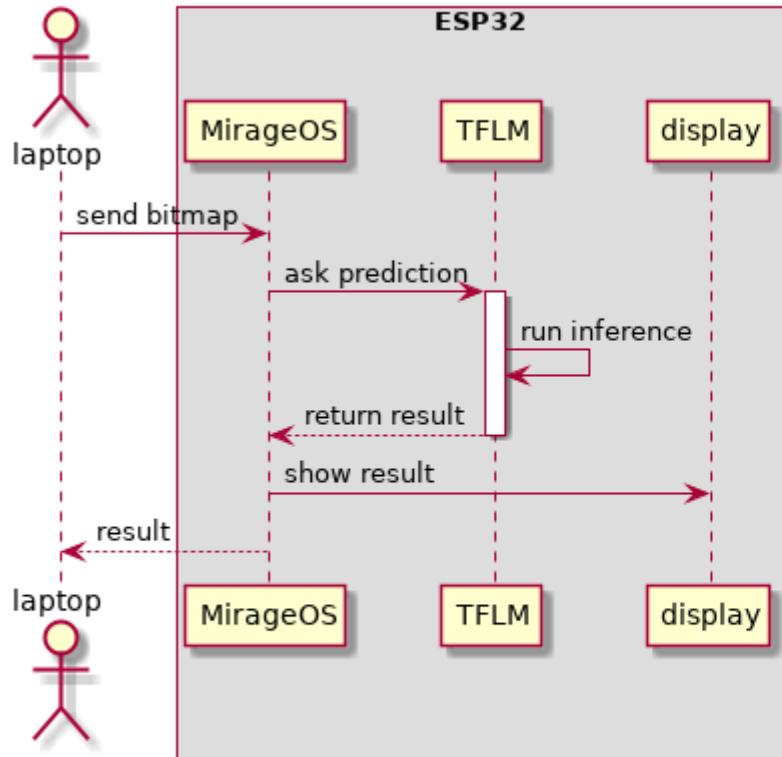
ML compiler: computation graph conversion



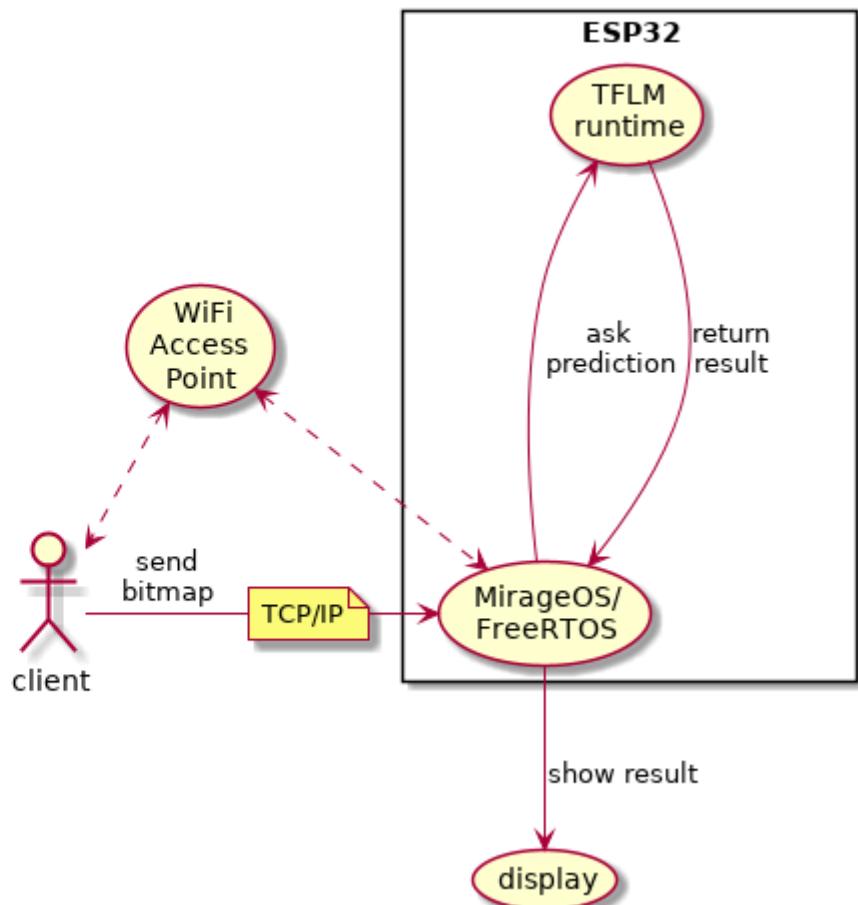
Outline

1. ~~Demo~~
2. ~~Problems~~
3. ~~Proposal~~
4. ~~Three Enablers~~
5. PoC
 - o ~~Training~~
 - o ~~Compile~~
 - o ~~Inference~~
6. Conclusion

Inference Sequence



Inference Usecase



Binary Size

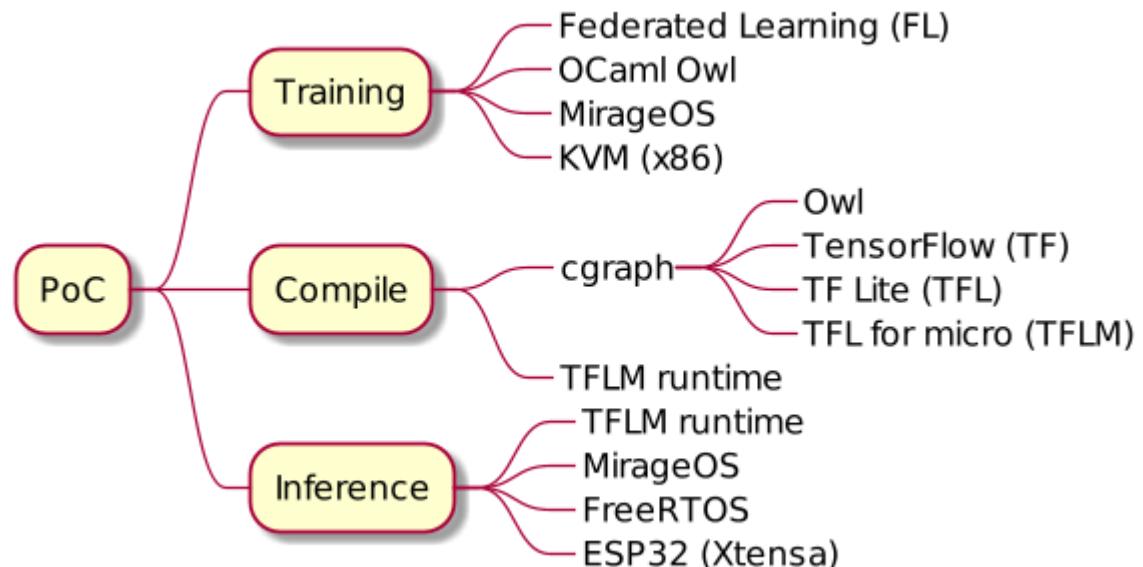
```
Total sizes: [35/93]
DRAM .data size: 18784 bytes
DRAM .bss size: 62936 bytes
Used static DRAM: 81720 bytes ( 99016 available, 45.2% used)
Used static IRAM: 91189 bytes ( 39883 available, 69.6% used)
    Flash code: 1062913 bytes
    Flash rodata: 854548 bytes
Total image size: ~2027434 bytes (.bin may be padded larger)
Per-archive contributions to ELF file:
          Archive File DRAM .data & .bss      IRAM Flash code & rodata   Total
            libmain.a      3410    40192        0     448568     300296  792466
            libtflm.a      1216        0        0     151465     447123  599804
            libnet80211.a    314     9056     3565    107308     12083  132326
  libc-psram-workaround.a    495       24     8070     83797      8017  100403
            liblwip.a       14     1962        0     61917     12227   76120
            libesp32.a      3105    2637    18693     21990     28550   74975
            libpp.a         1229    5274    13275     45188      4280   69246
<:- 3:- 4:-- 5:-*►64kB/s 25C 0.92 8x1.8GHz 31.3G18% 131.160.51.146 2019-09-
```

Outline

1. ~~Demo~~
2. ~~Problems~~
3. ~~Proposal~~
4. ~~Three Enablers~~
5. ~~PoC~~
6. **Conclusion**

Done

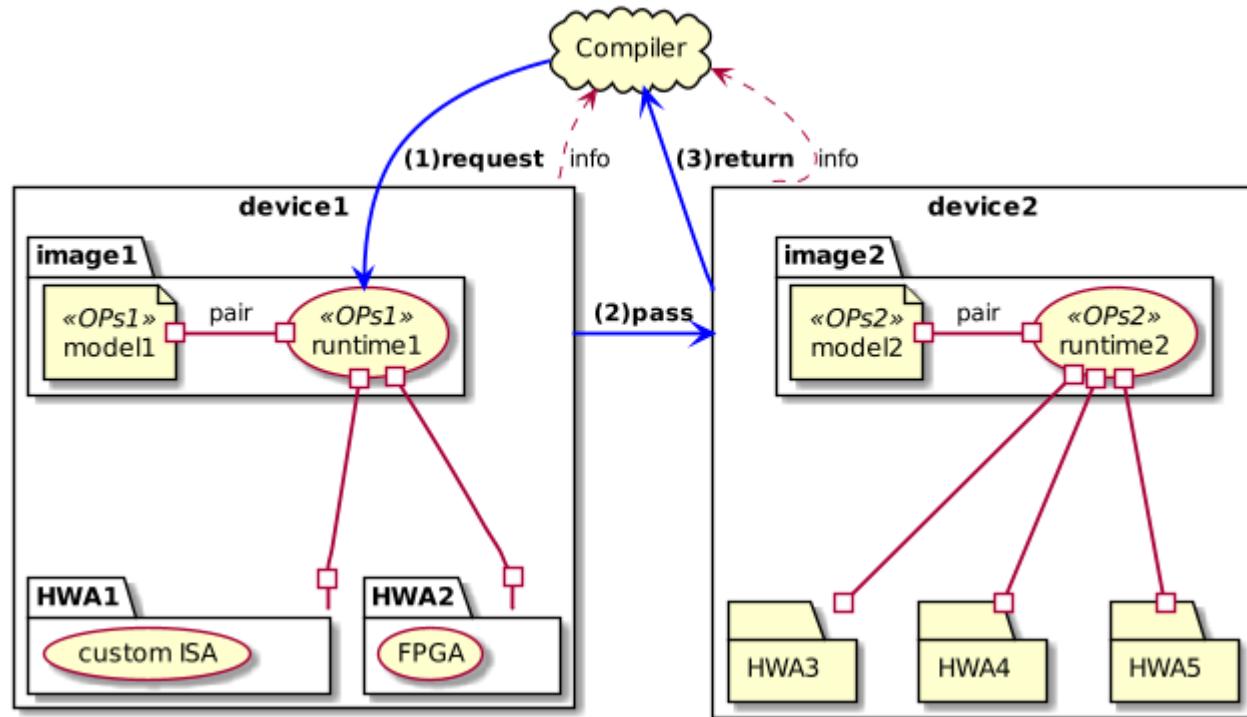
Demonstrated End-to-End TinyML (as-a-Service?)

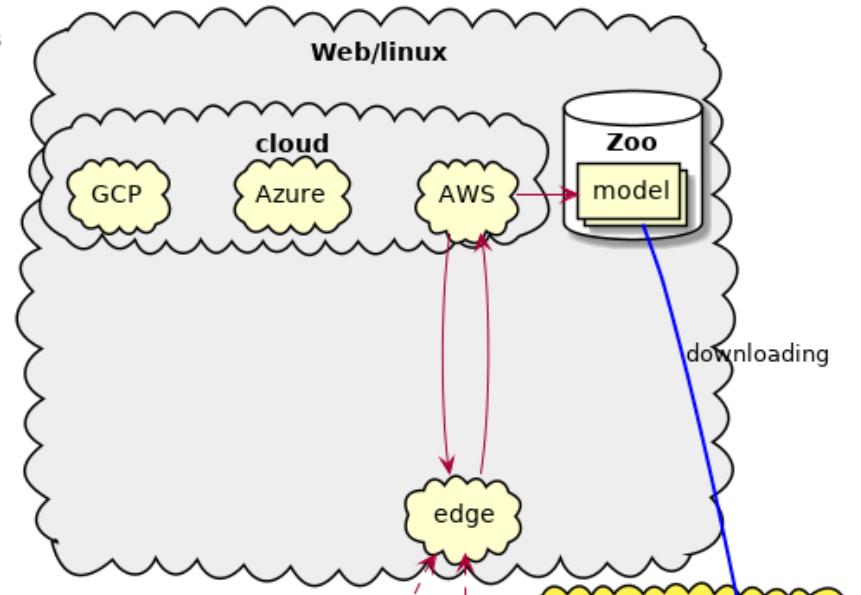
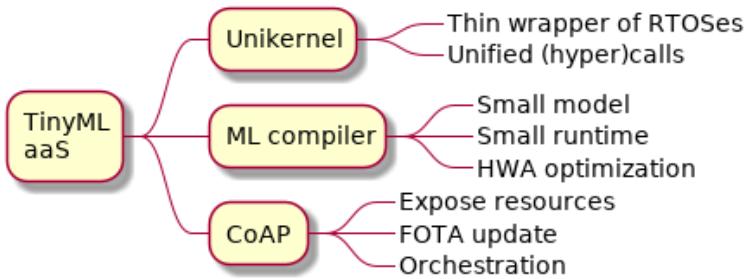


Next

1. Add **Orchestration** with CoAP?
2. Supprt other MCUs and/or **complicated** models?
3. **Heterogeneous** (distributed) training on MCU?
4. Distributed inference on **heterogeneous** HWAs?

multi-node, multi-HWA, scheduling

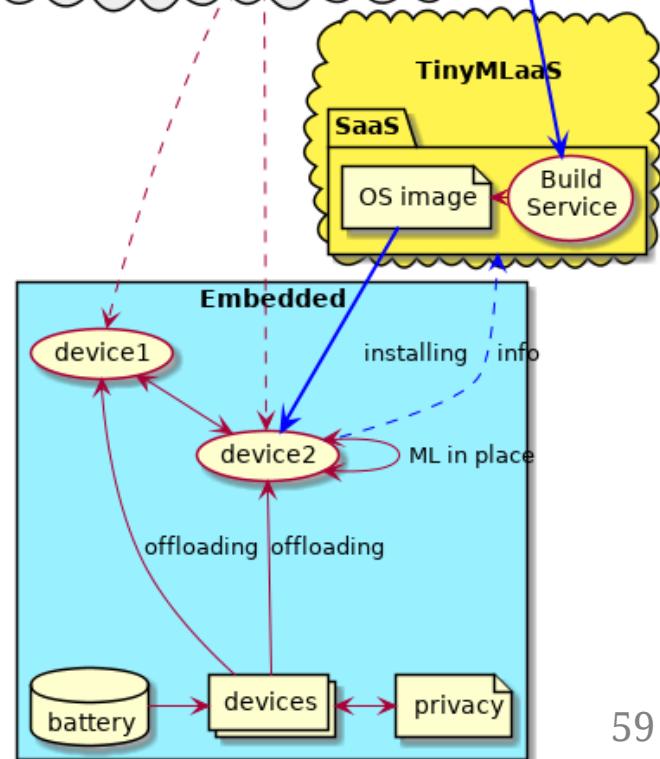




TinyML as-a-Service

could bring ML

onto IoT.

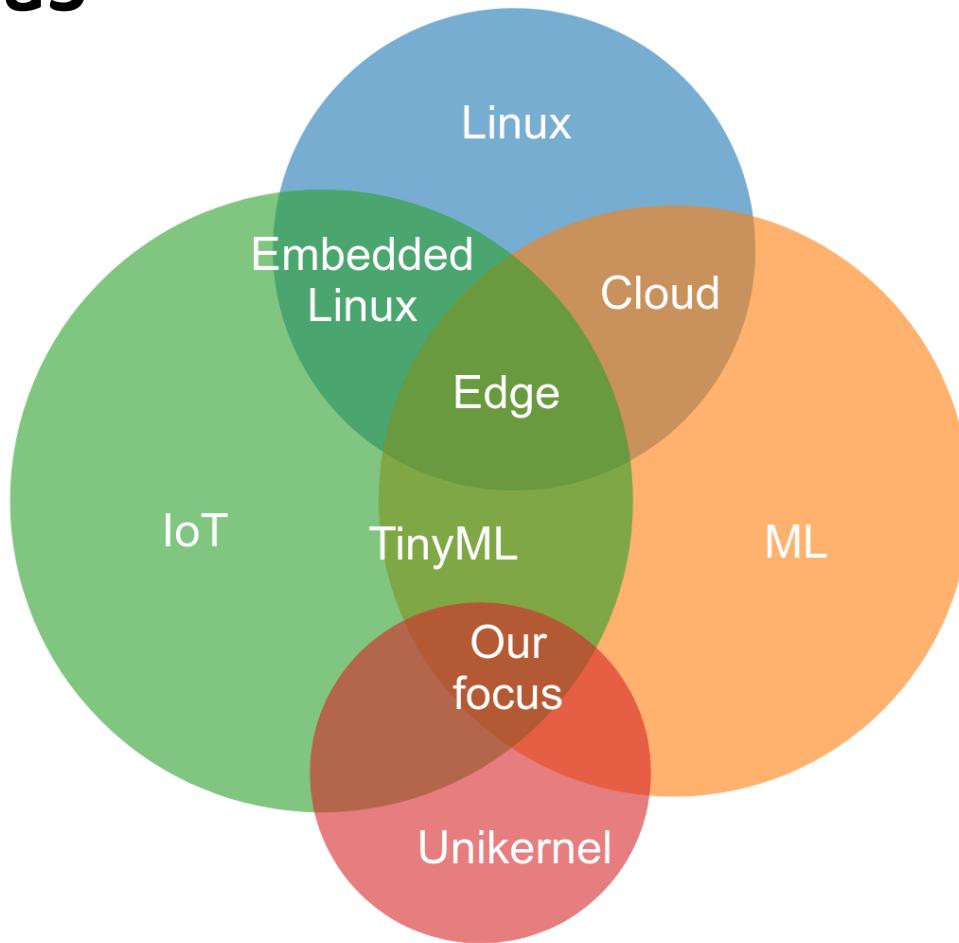




ERICSSON

Appendix

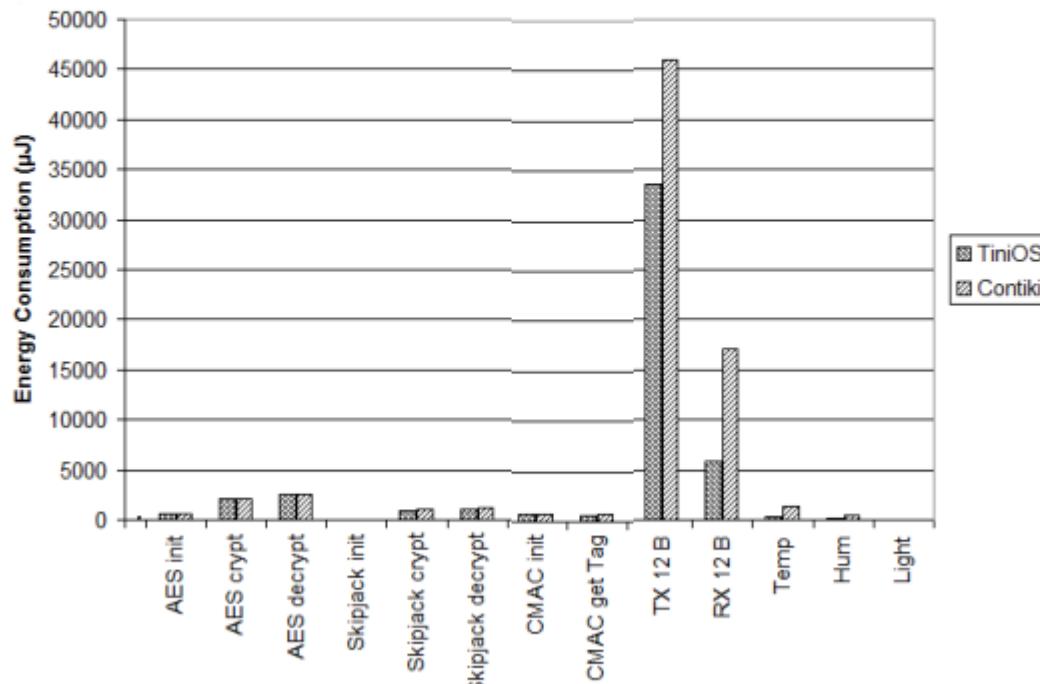
Our focus



Garrulity

Energy consumption on TelosB

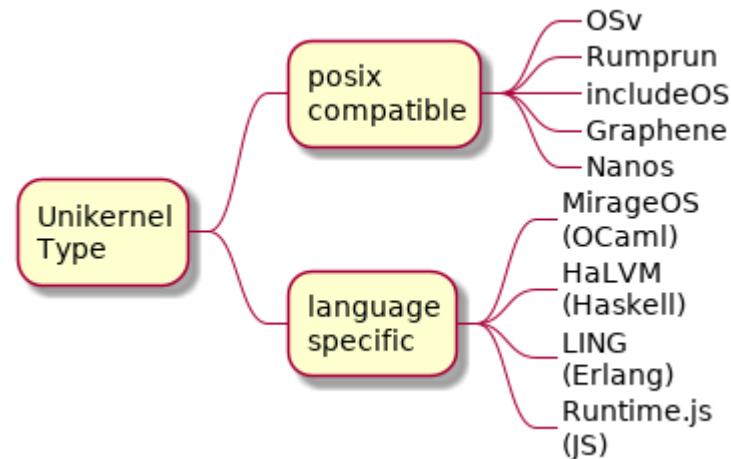
Message exchange cost: orders of magnitude more than processing, symmetric crypto



C.B. Margi, B.T. de Oliveira, G.T. de Souza, M.A. Simplicio Jr, P.S.L.M. Barreto,
T.C.M.B. Carvalho, M. Näslund, R. Gold, ICCCN'2010 / IEEE WiMAN 2010]

34

Type of Unikernel



MirageOS in OCaml

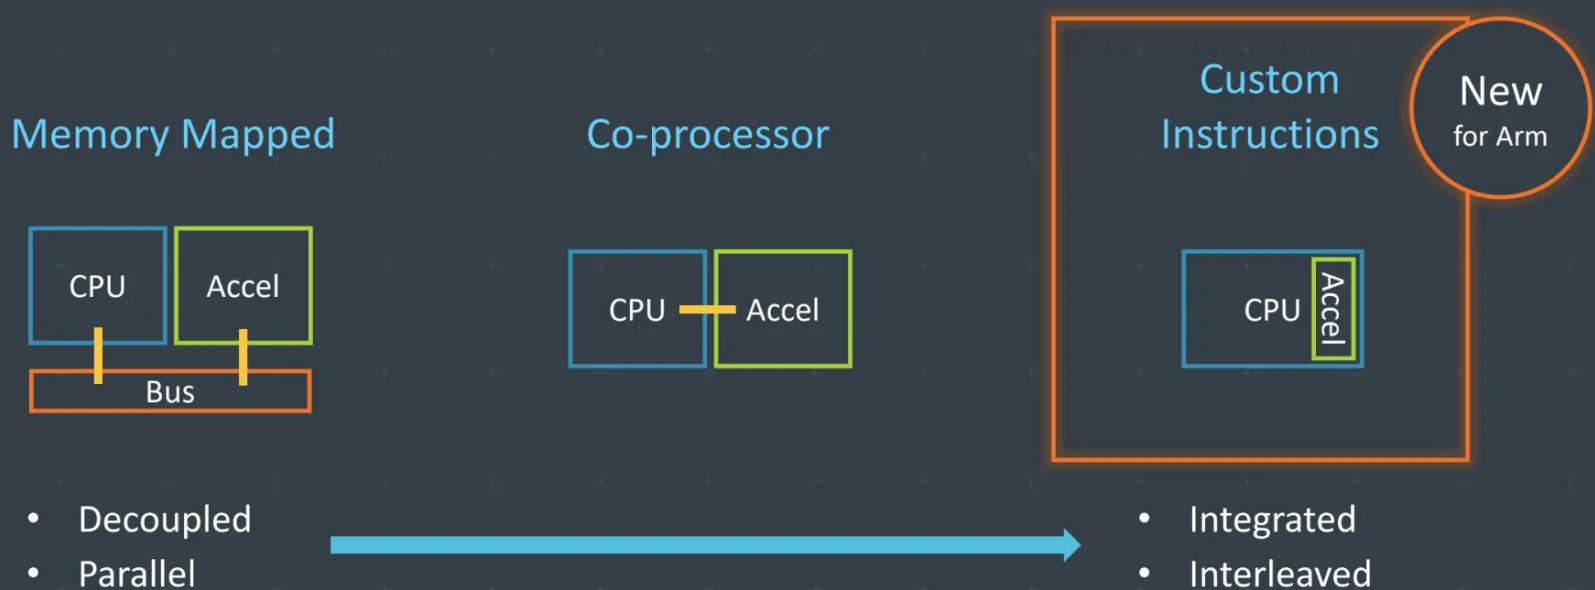
```
open Lwt.Infix

module Main (S: Mirage_stack_lwt.V4) = struct

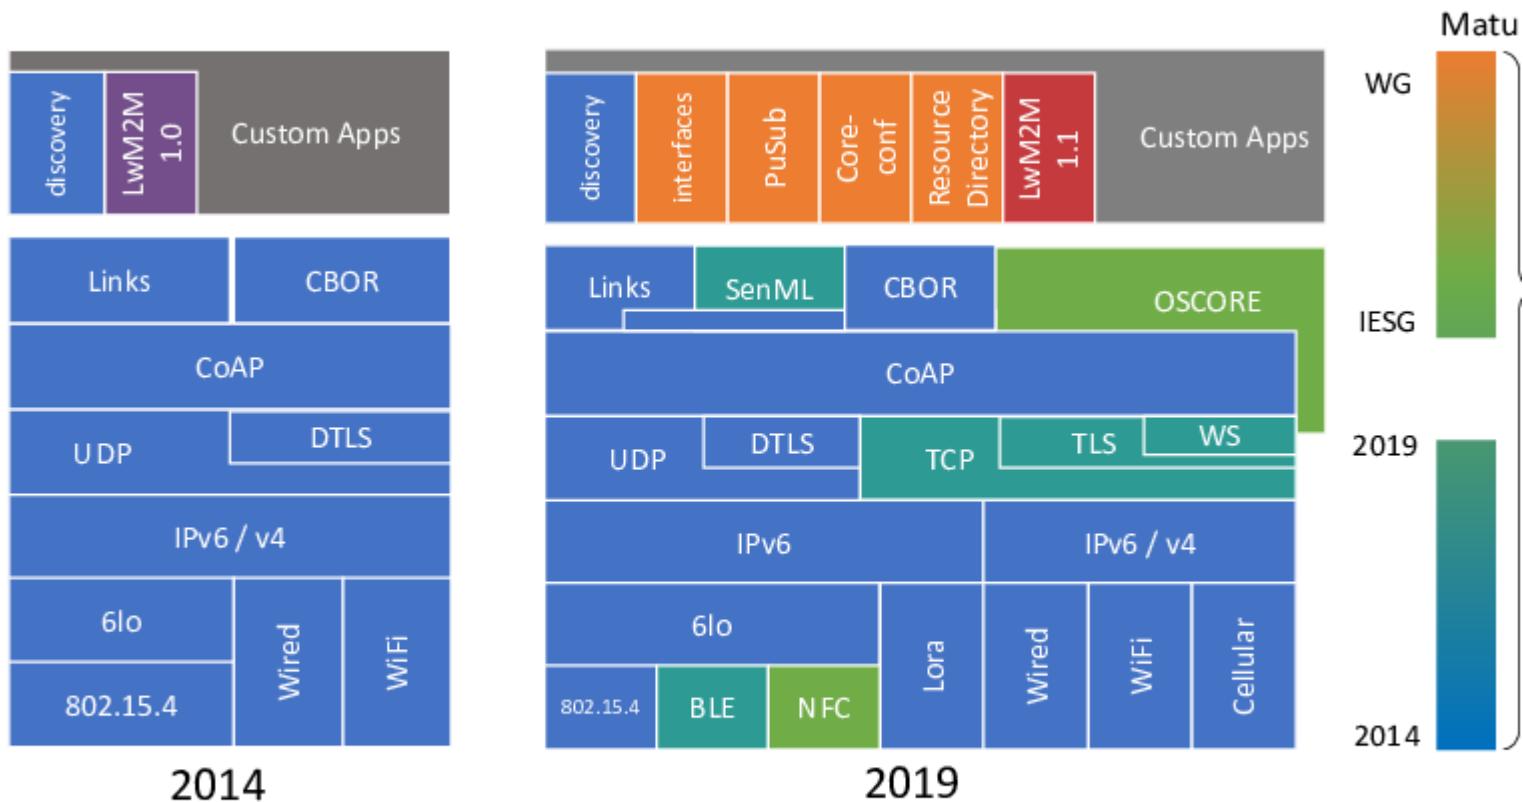
  let start s =
    let port = Key_gen.port () in
    S.listen_tcpv4 s ~port (fun flow ->
      let dst, dst_port = S.TCPV4.dst flow in
      Logs.info (fun f -> f "new tcp connection from IP %s on port %d$"
                  (Ipaddr.V4.to_string dst) dst_port);
      S.TCPV4.read flow >>= function
        | Ok `Eof -> Logs.info (fun f -> f "Closing connection!"); Lwt.$
        | Error e -> Logs.warn (fun f -> f "Error reading data from est$"
        | Ok (`Data b) ->
          Logs.debug (fun f -> f "read: %d bytes:\n%s" (Cstruct.len b) $
          S.TCPV4.close flow
      );
      S.listen s
    );
  end
-UU-:----F1 unikernel.ml All L1 Git-cdddd57 (Tuareg ARev Merlin (
Beginning of buffer
0:emacs#- 1:-* 2:>68kB/s 25C 0.59 8x1.2GHz 31.3G18% 131.160.51.146 2019
```

Custom ISA, RISC-V & Arm

Meeting the requirements for acceleration



Standards Device Stack

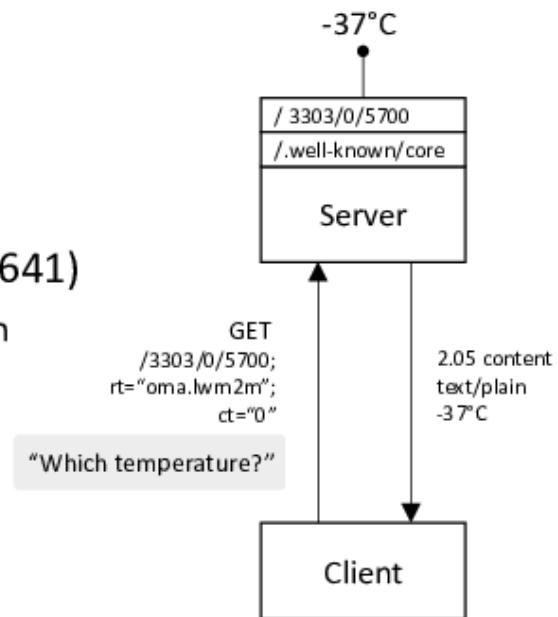


13/2/19

One Data Model Group

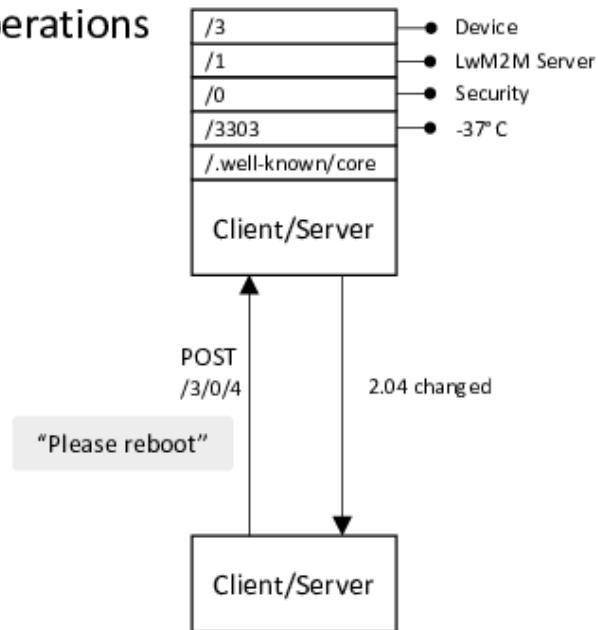
The Constrained Application Protocol (CoAP)

- CoRELink (RFC6690) provides a link format
 - Reuses Web Linking RFC5988 for IoT.
 - Enables query parameters for discovery (lt, gt...)
 - Enables attribute and relation types (rt, if, sz).
`<3303/0/5700>;rt="oma:lwm2m:temp";ct="0"`
- Notifications available through *observe* option (RFC7641)
 - Can observe and add query parameters to the observation
`<3303/0/5700?lt=0>`
- The “*/.well-known/core*” URI provides discovery
- Multiple serialization formats used with CoAP
 - SenML (RFC8428): Minimalistic JSON
 - CBOR (RFC7049): Binary serialization
- Multiple implementations available at coap.technology

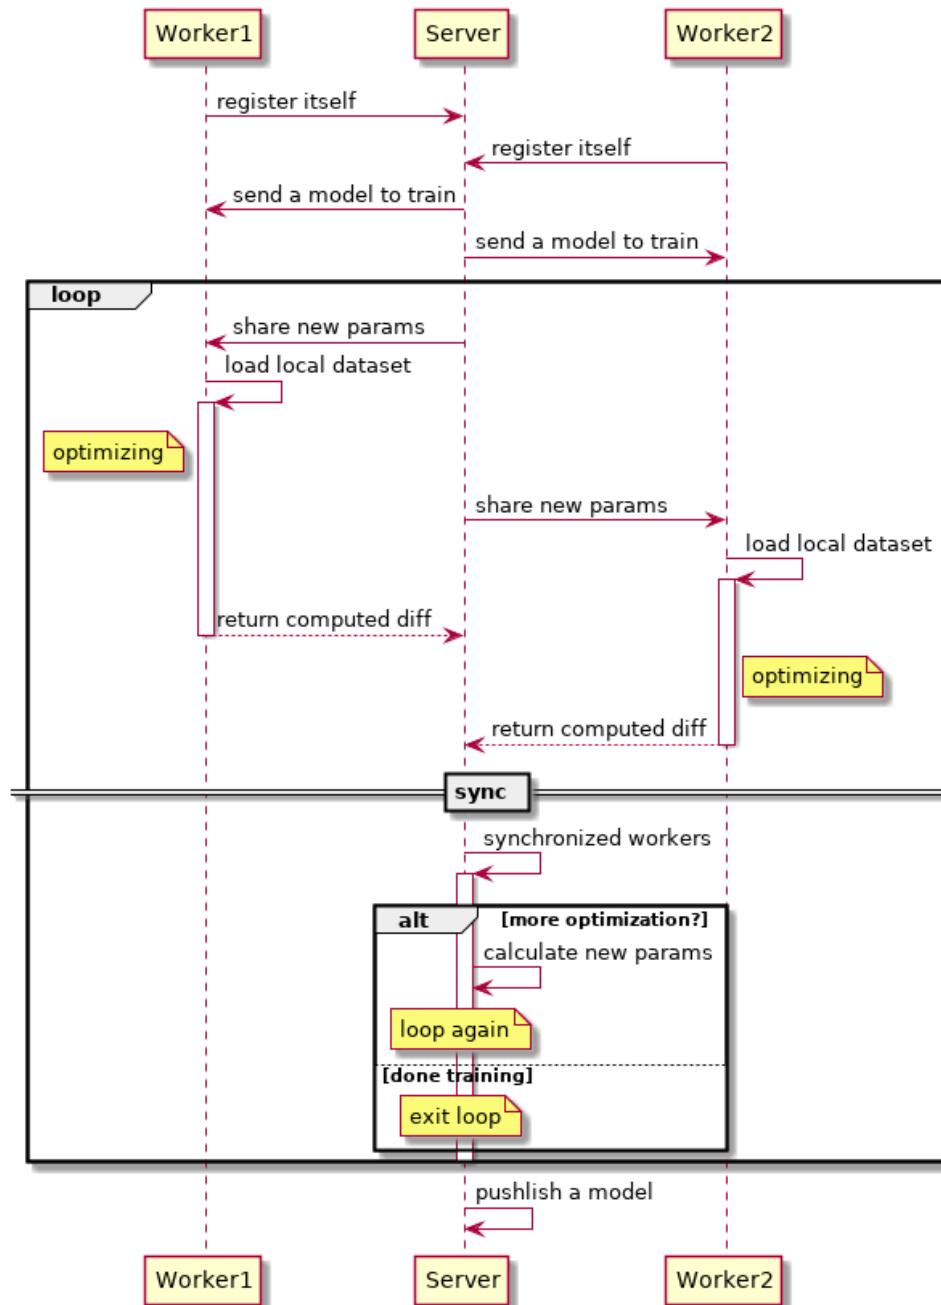


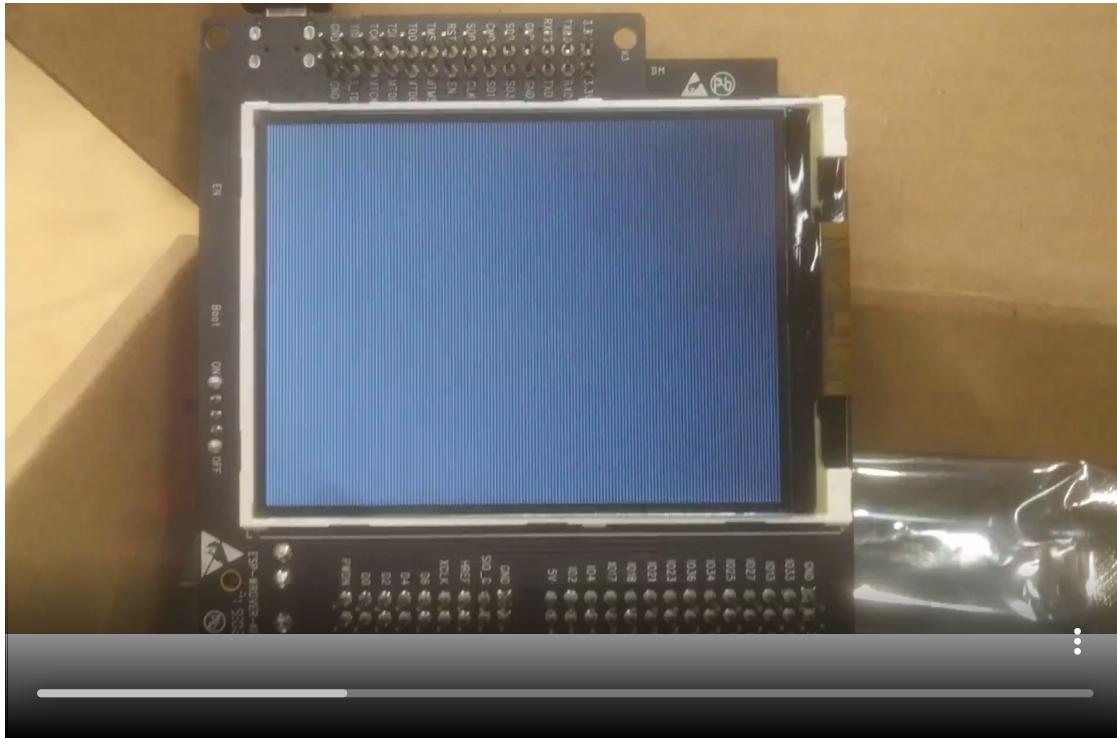
The LightWeight M2M Protocol (LwM2M)

- Mapping of CoAP methods (GET, POST, PUT...) to **CRUD** operations
- Interaction with device through simple “Objects”
 - RWX, Access Control, Observation, Notification
 - Independent from underlying protocol stack (CoAP today)
 - Simple resource structure
 - Objects’ resources are accessed with simple URIs:
$$\{/Object\ ID\}/\{Object\ Instance\}/\{Resource\ ID\}$$
 - Multiple serializations:
For example JSON, CBOR and raw values.
- Common repository for all Objects (OMNA)
 - Enables interoperability and reusability



Training: Parameter Server(PS) sequence





Fashion MNIST from Zalando